

Deep Learning for Computer Vision

( Spring 2018

<http://vllab.ee.ntu.edu.tw/dlcv.html> (primary)

<https://ceiba.ntu.edu.tw/1062DLCV> (grade, etc.)

FB: [DLCV Spring 2018](#)

Yu-Chiang Frank Wang 王鈺強, Associate Professor
Dept. Electrical Engineering, National Taiwan University

About BONUS in HW1

- “Given a $d \times d$ symmetric matrix”
→ “Given a $d \times d$ **positive semi-definite** symmetric matrix”
 - We still assume that this matrix has distinct eigenvalues.
 - ~~☆~~ And, you need to explain why adding this statement matters.
-
- Deadline for BONUS submission (3/23 Friday 1:00am)
 - No late day for the BONUS.



TA Hours

Mon	Tue	Wed	Thu	Fri
 Yu-Jhe Li 11:00-12:00	 Yan-Bo Lin 14:20-15:10	 Alex Liu 13:20-14:20	 Chi-Hsin Lo 10:00-11:00	 曾耕森 13:20-14:20
 Huai-Jin Peng 13:20-14:20	 Hsuan-I Ho 15:30-17:20	 楊繼程 14:30-15:30	 Cheng-Yen Yang 13:30-14:30	 Jia-Wei Yan 15:10-16:10
 Frank Liao 18:30-19:30	 Jason Kuo 18:00-19:00	 Yu-Ying Yeh 14:30-15:30	 Boris Chen 15:30-16:30	

- FB: [DLCV Spring 2018](#)
- Feedback are welcome (to TA/me).

About Grading & Final Project

- HW assignments (60%)
 - HW 1 (5%), HW 2 (10%), HWs 3~5 (15% each)
 - Final project + report + poster presentation (35%)
 - Participation (5%)
 - Mostly 課堂/課後/facebook群組互動、Q&A等
-
- About Final Project
 - **3 (preferable) or 4 (max) people per group**; plz team up early!
 - Details (e.g., computing resource, topics, prize, etc.): TBA in April
 - **Bonus Points (5%)**
 - Turn your final project/report with **sufficient quality** into contests or major/top-tier conference submissions.

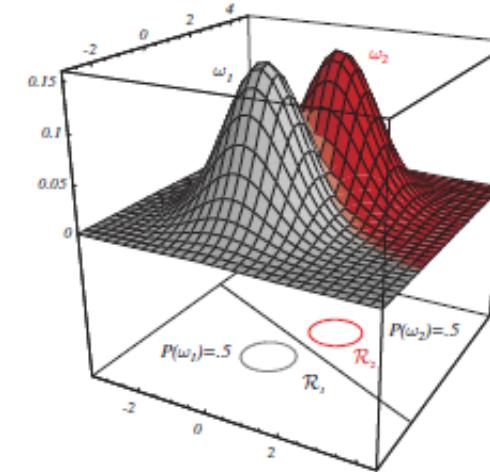
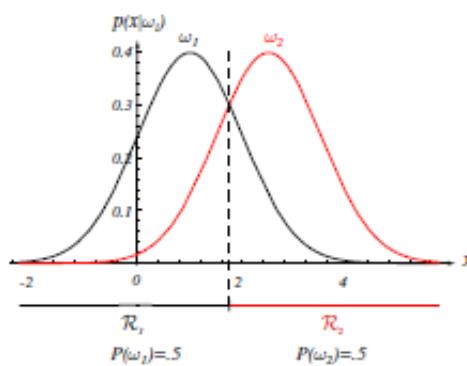


Tight (yet tentative) Schedule

Week	Date	Topic	Remarks
0	3/7	Course Logistics + Intro to Computer Vision	
1	3/14	Machine Learning 101	HW #1 out (due in 1 week)
2	3/21	Image Representation	HW #1 due (Bonus due 3/23 Fri. 1am)
3	3/28	Interest Point: From Recognition to Tracking	HW #2 out
4	4/4	N/A	Spring Break
5	4/11	Intro to Neural Networks + CNN	HW #2 due
6	4/18	Detection & Segmentation	HW #3 out
7	4/25	Generative Models	
8	5/2	Visualization and Understanding NNs	HW #3 due / HW #4 out
9	5/9	Recurrent NNs and Seq-to-Seq Models (I)	Project team
10	5/16	Recurrent NNs and Seq-to-Seq Models (II)	HW #4 due / HW #5 out
11	5/23	Deep Reinforcement Learning for Visual Apps	Project proposal
12	5/30	Learning Beyond Images (2D/3D, depth, etc.)	
13	6/6	Transfer Learning for Visual Analysis	HW #5 due
14	6/13	Vision & Language / Guest Lectures	Project midterm progress
15	6/20	Industrial Visit (e.g., Microsoft @ TPE)	CVPR week
16	6/27 (?)	Final Project Presentation	

What We Covered Last Week...

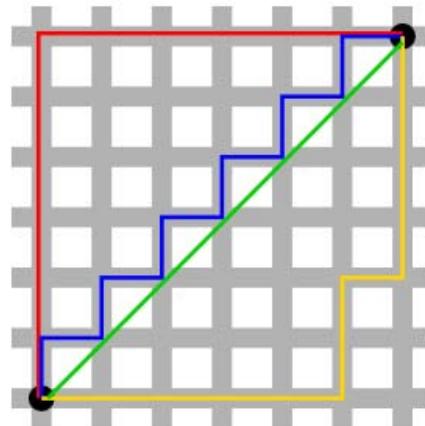
- From Probability to Bayes Decision Rule
- Brief Review of Linear Algebra & Linear System
- Unsupervised vs. Supervised Learning
 - Clustering & Dimension Reduction
 - Training, testing, & validation
 - Linear Classification



Revisit of Clustering

- Similarity:
 - A key component/measure to perform data clustering
 - Inversely proportional to distance
 - Example distance metrics:

- Euclidean distance (L2 norm): $d(x, z) = \|x - z\|_2 = \sqrt{\sum_{i=1}^D (x_i - z_i)^2}$
- Manhattan distance (L1 norm): $d(x, z) = \|x - z\|_1 = \sum_{i=1}^D |x_i - z_i|$



Revisit of Clustering (cont'd)

- Similarity:

- A key component/measure to perform data clustering

- Inversely proportional to distance

- Example distance metrics:

- Kernelized (non-linear) distance:

$$d(x, z) = \|\Phi(x) - \Phi(z)\|_2^2 = \|\Phi(x)\|_2^2 + \|\Phi(z)\|_2^2 - 2\Phi(x)^T \Phi(z)$$

- No need to know explicit Φ , just the way to calculate $\Phi(x)^T \Phi(z)$.

- E.g., Gaussian kernel: $K(x, z) = \exp\left(-\frac{\|x - z\|_2^2}{2\sigma^2}\right)$, we have

$$\Phi(x)^T \Phi(z) = \Phi(x)^T \Phi(z) = 1 \quad \text{if } \|x - z\|_2^2 < \Phi(x)^T \Phi(z) < 1$$

- Distance is more sensitive to larger/smaller σ ?

- E.g., Given the same data pair x and z , $\Phi(x)^T \Phi(z) \rightarrow 0$ if a very large/small σ .

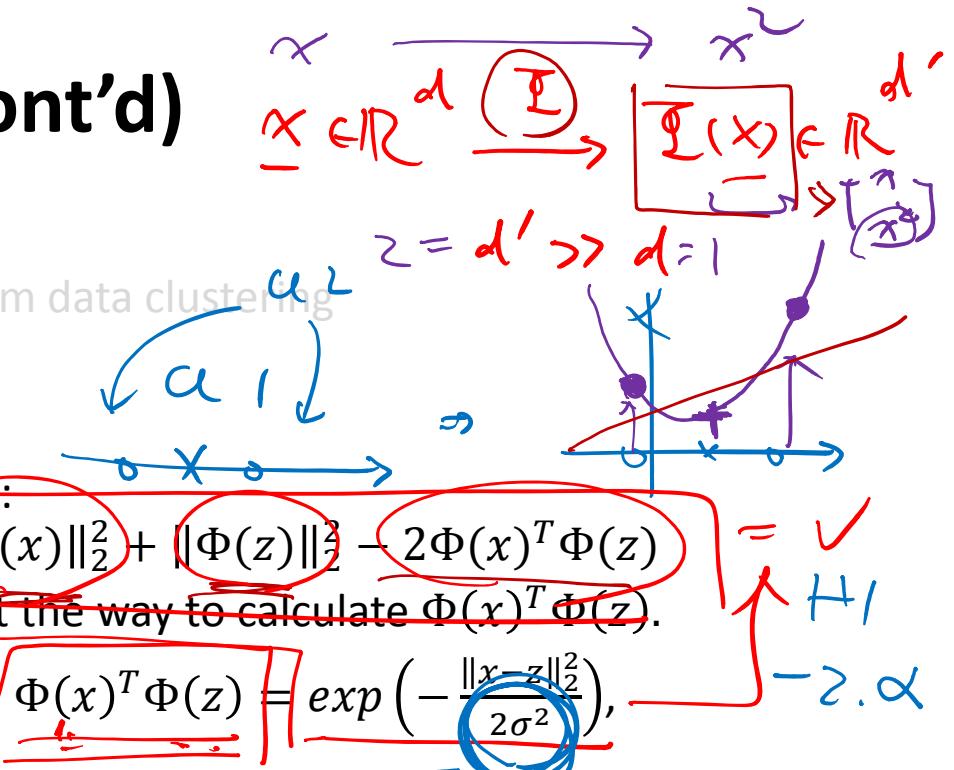
$$x \leftrightarrow z$$

$\sigma \uparrow$

$\sigma \downarrow$

$$\exp\left(\frac{0}{\sigma^2}\right) \rightarrow K \rightarrow 1$$

$$\exp\left(\frac{0}{\sigma^2}\right) \rightarrow K \rightarrow 0$$



K-Means Clustering

- **Input:** N examples $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ ($\mathbf{x}_n \in \mathbb{R}^D$); number of partitions K
- **Initialize:** K cluster centers μ_1, \dots, μ_K . Several initialization options:
 - Randomly initialize μ_1, \dots, μ_K anywhere in \mathbb{R}^D
 - Or, simply choose any K examples as the cluster centers
- **Iterate:**
 - Assign each of example \mathbf{x}_n to its closest cluster center
 - Recompute the new cluster centers μ_k (mean/centroid of the set C_k)
 - Repeat while not converge
- **Possible convergence criteria:**
 - Cluster centers do not change anymore
 - Max. number of iterations reached
- **Output:**
 - K clusters (with centers/means of each cluster)

K-Means Clustering (cont'd)

- Proof in 1-D case (if time permits). (X)

K-Means Clustering (cont'd)

- Remarks
 - Standard K-means clustering
 - Hard assignment only.
 - Mathematically, we have:

$$\min_{\underline{A}, \underline{\mathcal{Z}}_n} \sum_{i=1}^N \| \underline{x}_i - [\underline{A}] \underline{\mathcal{Z}}_i \|_2^2$$

$\uparrow \quad \quad \quad \uparrow \quad \quad \quad \uparrow$
 $d \times 1 \quad d \times k \quad k \times 1$

$\underline{\mathcal{Z}}_i = \begin{bmatrix} \dots \\ \vdots \end{bmatrix}$
 $\underline{\mu}_1 \quad \underline{\mu}_k$

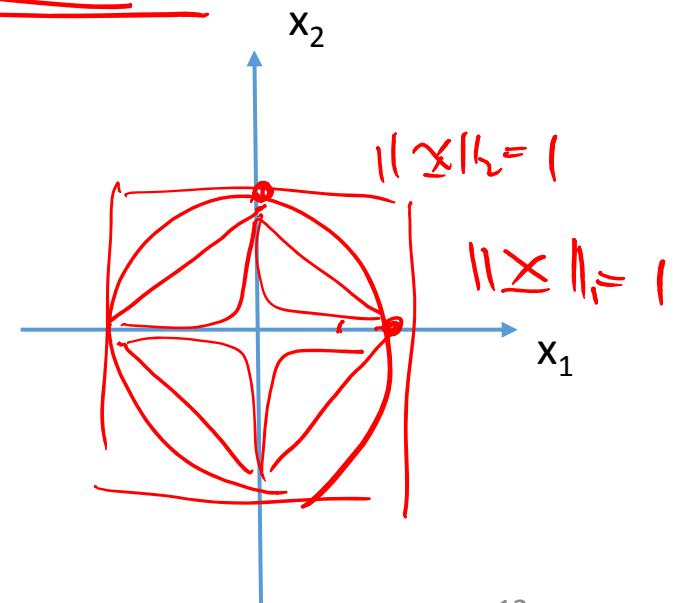
$\| \underline{x}_i - \underline{\mu}_1 \|_2^2$
 \rightarrow hard

$\left\{ \begin{array}{l} \| \underline{\mathcal{Z}}_i \|_0 = 1 \\ \| \underline{\mathcal{Z}}_i \|_1 = 1 \end{array} \right.$
 $\uparrow \quad \quad \quad \uparrow$
 \rightarrow soft

K-Means Clustering (cont'd)

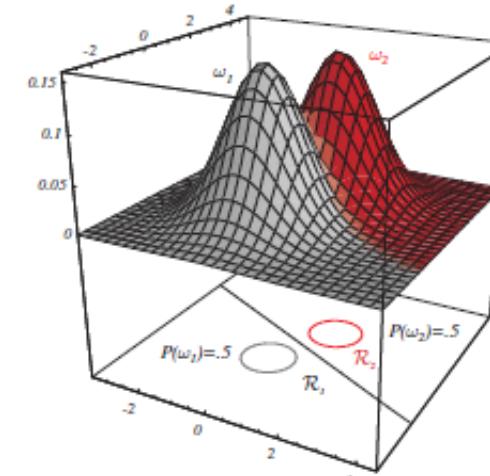
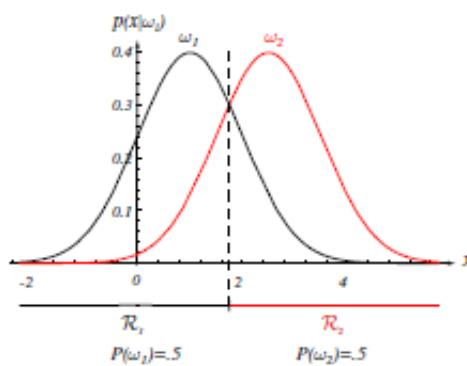
- Remarks

- Recall that, we have $\|x\|_p = \left(\sum_{i=1}^d x_i^p \right)^{1/p}$
- L₂/Euclidean norm $\|x\|_2 = \sqrt{\sum_{i=1}^d x_i^2}$
- L₁ norm $\|x\|_1 = \sum_{j=1}^d |x_j| \rightarrow 1$
- Maximum norm $\|x\|_\infty = \max\{|x_1|, \dots, |x_d|\} = 1$
- L₀ norm: $\|x\|_0 = \lim_{p \rightarrow 0} \left(\sum_{i=1}^d x_i^p \right)^{1/p}$
 $= (\underbrace{x_1 + \dots + x_d}_{\# \text{ of non-zero elements}})^{1/p}$



What We Covered Last Week...

- From Probability to Bayes Decision Rule
- Brief Review of Linear Algebra & Linear System
- Unsupervised vs. Supervised Learning
 - Clustering & Dimension Reduction
 - Training, testing, & validation
 - Linear Classification



PCA vs. SVD

- Formulation of PCA and Data Whitening

$$\Sigma = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)(x_i - \mu)^T$$

$$\rightarrow \Sigma = E \cdot \Delta \cdot E^T$$

$$E = [e_1 | \dots | e_d]$$

$$e_1 + \dots + e_d = d$$

$$= \bar{x} \cdot \bar{x}^T = d \times d$$

$$\begin{bmatrix} x_1 \\ \vdots \\ x_N \end{bmatrix} - \mu = \begin{bmatrix} x_1 - \mu \\ \vdots \\ x_N - \mu \end{bmatrix}$$

$$\Sigma = E \cdot \Delta \cdot E^T$$

$$\Sigma = E \cdot \Delta \cdot E^T$$

$$\Sigma \cdot e_i = \lambda_i \cdot e_i$$

$$d \times d \quad d \times 1$$

$$\Sigma = E \cdot \Delta \cdot E^T = \begin{bmatrix} e_1 & \dots & e_d \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_d \end{bmatrix} \begin{bmatrix} e_1^T \\ \vdots \\ e_d^T \end{bmatrix}$$

PCA vs. SVD (cont'd)

- Formulation of SVD & Its Relation to PCA

The diagram illustrates the relationship between Principal Component Analysis (PCA) and Singular Value Decomposition (SVD). It shows the decomposition of a matrix \underline{X} into $\underline{U} \cdot \underline{\Sigma} \cdot \underline{V}^T$. The matrix \underline{X} is $d \times N$. The matrix $\underline{\Sigma}$ is $d \times d$, \underline{U} is $d \times d$, and \underline{V}^T is $N \times d$.

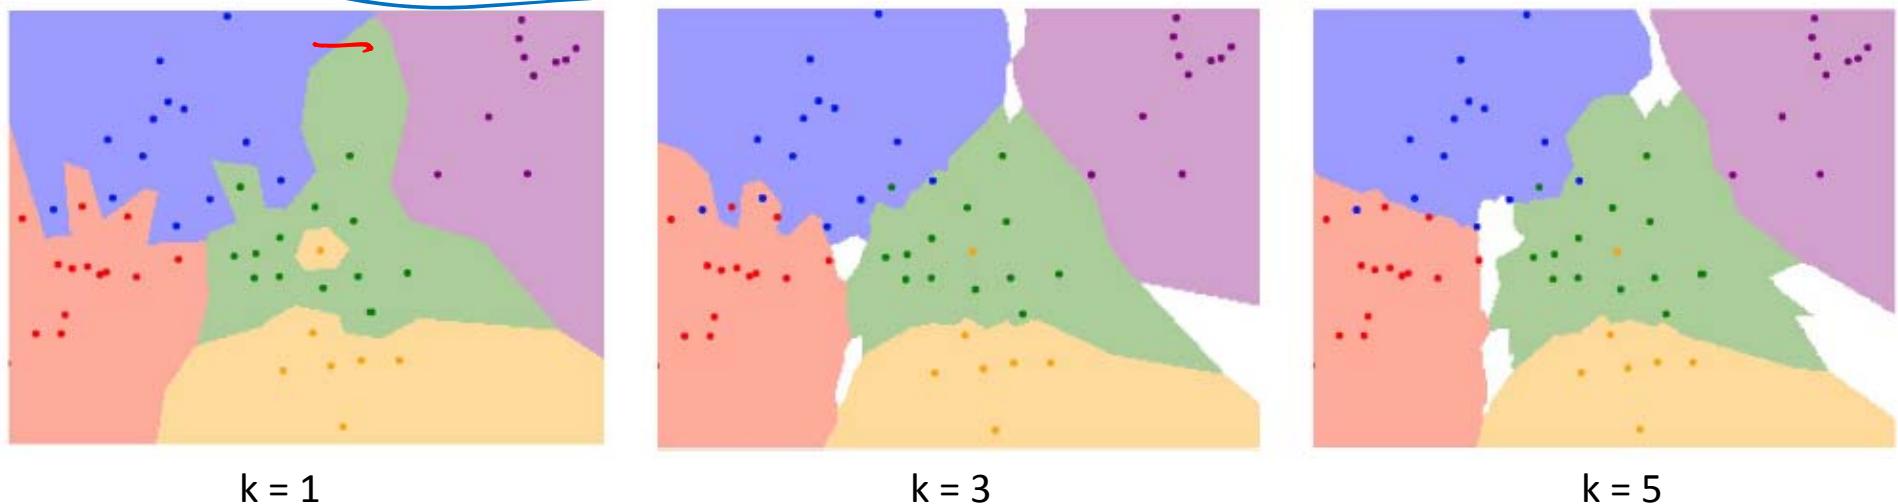
On the right, the matrix $\underline{\Sigma}$ is shown as a diagonal matrix with entries $\sigma_1, \sigma_2, \dots, \sigma_d$. Above it, the matrix $\underline{\Sigma} = \underline{U}^T \cdot \underline{\Sigma} \cdot \underline{U}$ is shown, where \underline{U}^T is $N \times d$ and $\underline{\Sigma}$ is $d \times d$.

Below, the equation $\underline{X} \cdot \underline{X}^T = \underline{U} \cdot \underline{\Sigma} \cdot \underline{V}^T (\underline{V} \cdot \underline{\Sigma}^T \cdot \underline{U}^T)$ is derived.

At the bottom, the matrix $\underline{\Sigma}$ is shown as $\underline{\Sigma} = \underline{U} \cdot \underline{\Lambda} \cdot \underline{U}^T$, where $\underline{\Lambda}$ is a diagonal matrix with entries $\sigma_1^2, \sigma_2^2, \dots, \sigma_d^2$.

Remarks on Nearest-Neighbor Classification

- How do we classify? For example...
 - Given a test face input, project into the same 3D space (by the same 3 eigenvectors).
 - The resulting vector in the 3D space is the **feature** for this test input.
 - We can do a simple **Nearest Neighbor (NN)** classification with Euclidean distance, which calculates the distance to all the **projected training data** in this space.
 - If NN, then the **label of the closest training instance** determines the classification output.
 - If **k-nearest neighbors (k-NN)**, then k-nearest neighbors need to **vote** for the decision.

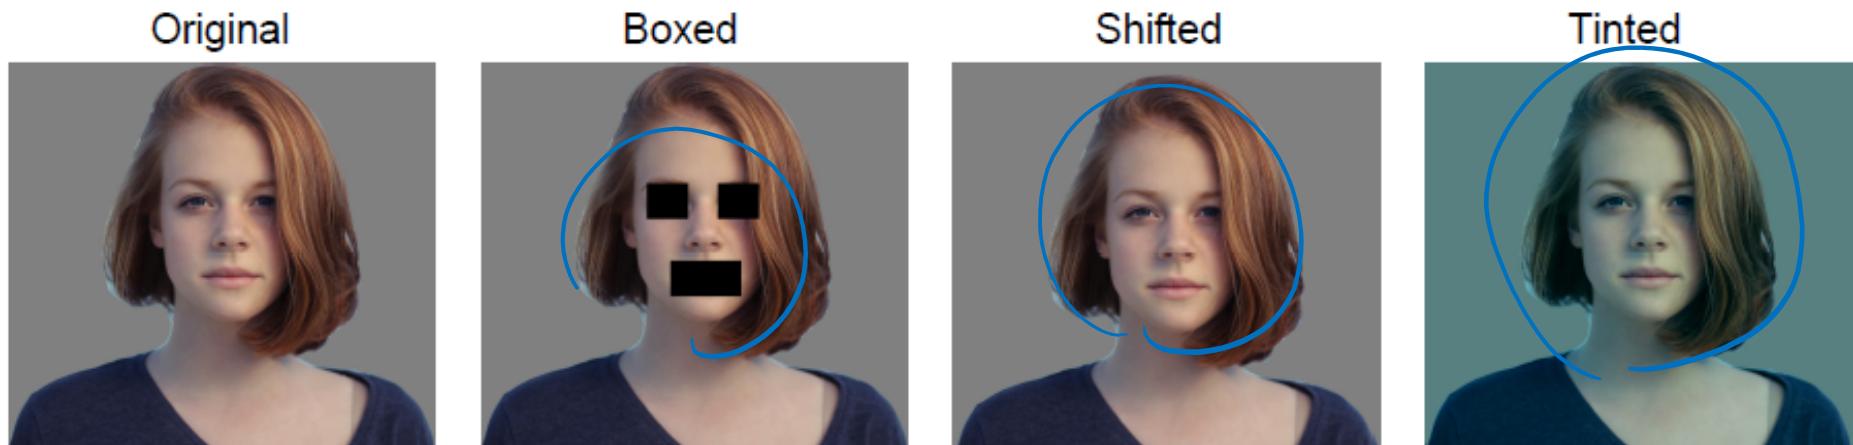


Demo available at <http://vision.stanford.edu/teaching/cs231n-demos/knn/>

Image credit: Stanford CS231n

Minor Remarks on NN-based Methods

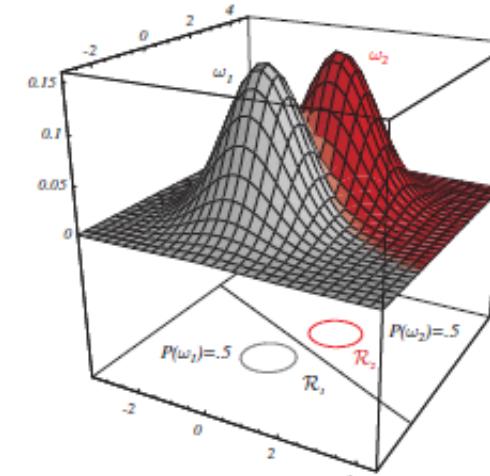
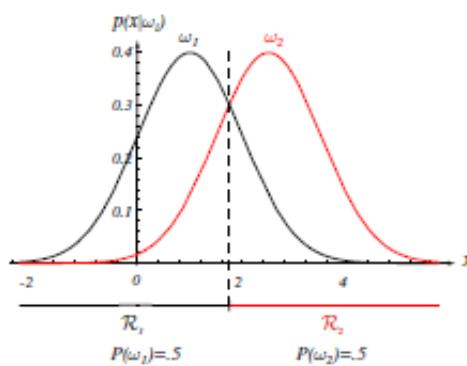
- In fact, k-NN (or even NN) is not of much interest in practice. Why?
 - Choice of *distance metrics* might be an issue. See example below.
 - Measuring distances in *high-dimensional spaces* might not be a good idea.
 - Moreover, NN-based methods require lots of *calculation* and *storage*!
(That is why NN-based methods are viewed as *data-driven* approaches)



All three images have the same Euclidean distance to the original one.

What We Covered Last Week...

- From Probability to Bayes Decision Rule
- Brief Review of Linear Algebra & Linear System
- Unsupervised vs. Supervised Learning
 - Clustering & Dimension Reduction
 - Training, testing, & validation
 - Linear Classification



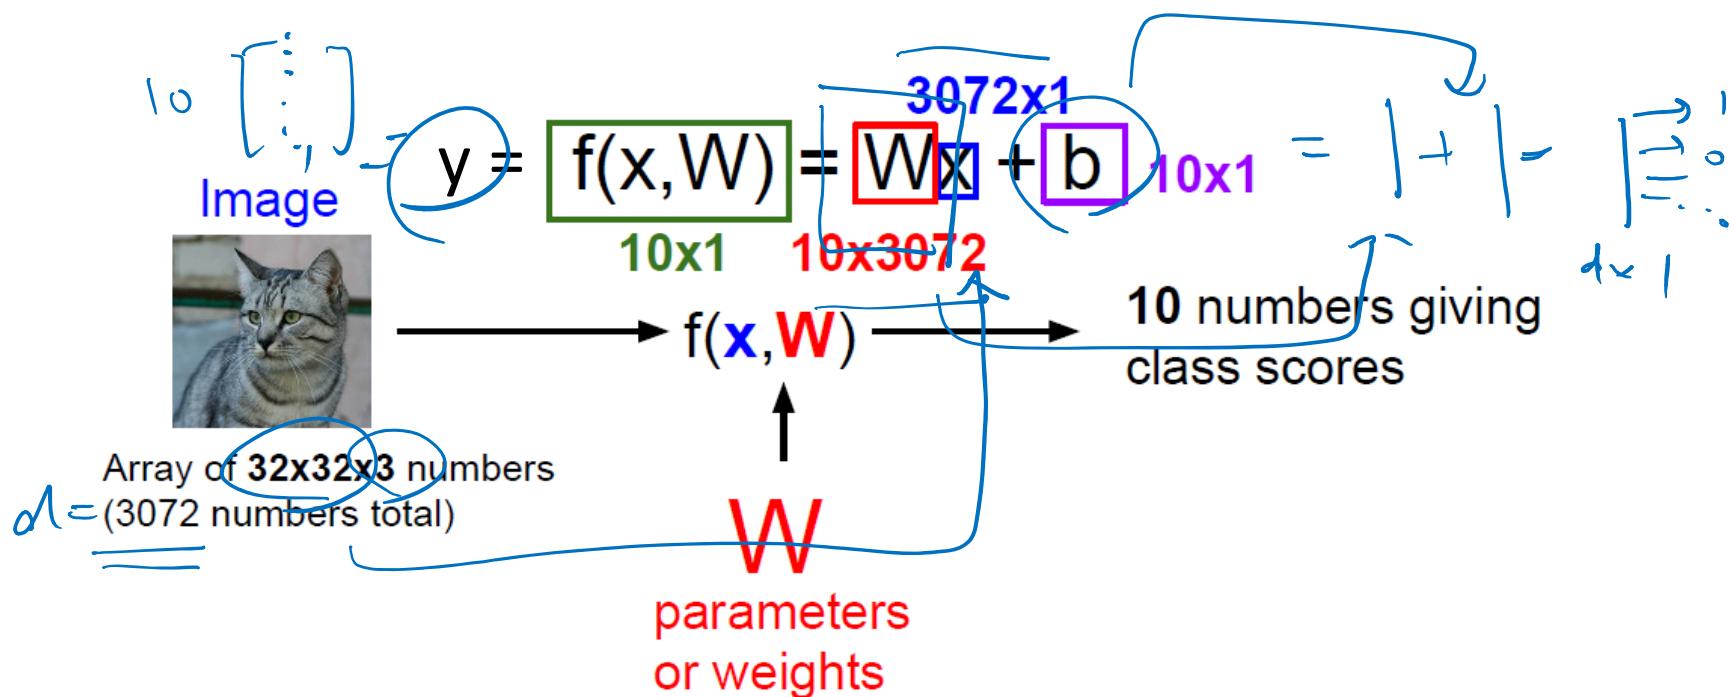
Linear Classification

- Linear Classifier
 - Can be viewed as a **parametric approach**, as discussed later.
 - Assuming that we need to recognize 10 object categories of interest
 - E.g., CIFAR10 with 50K training & 10K test images of 10 categories.
And, each image is of size 32 x 32 x 3 pixels.



Linear Classification (cont'd)

- Linear Classifier
 - Can be viewed as a **parametric approach**. Why?
 - Assuming that we need to recognize 10 object categories of interest (e.g., CIFAR10).
 - Let's take the input image as \mathbf{x} , and the linear classifier as \mathbf{W} . We hope to see that $\mathbf{y} = \mathbf{Wx} + \mathbf{b}$ as a 10-dimensional output indicating the score for each class.



Linear Classification (cont'd)

- Linear Classifier
 - Can be viewed as a **parametric approach**. Why?
 - Assuming that we need to recognize 10 object categories of interest (e.g., CIFAR10).
 - Let's take the input image as \mathbf{x} , and the linear classifier as \mathbf{W} . We hope to see that $\mathbf{y} = \mathbf{Wx} + \mathbf{b}$ as a **10-dimensional output** indicating the score for each class.
 - Take an image with 2×2 pixels & 3 classes of interest as example: we need to learn linear transformation/classifier \mathbf{W} and bias \mathbf{b} , so that desirable outputs $\mathbf{y} = \mathbf{Wx} + \mathbf{b}$ can be expected.

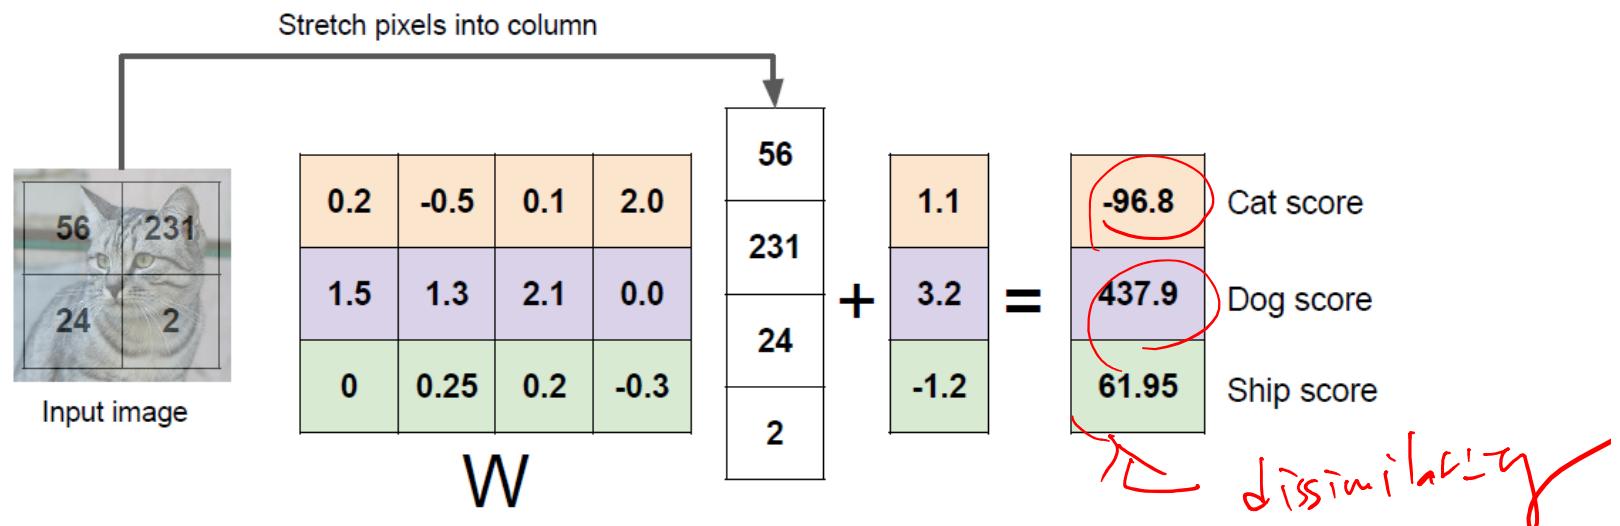
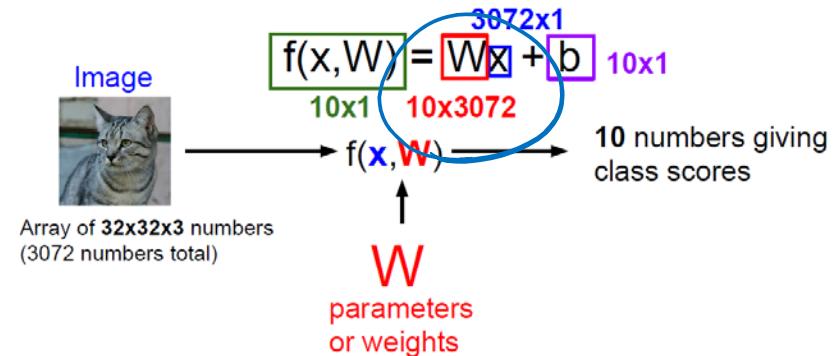


Image credit: Stanford CS231n

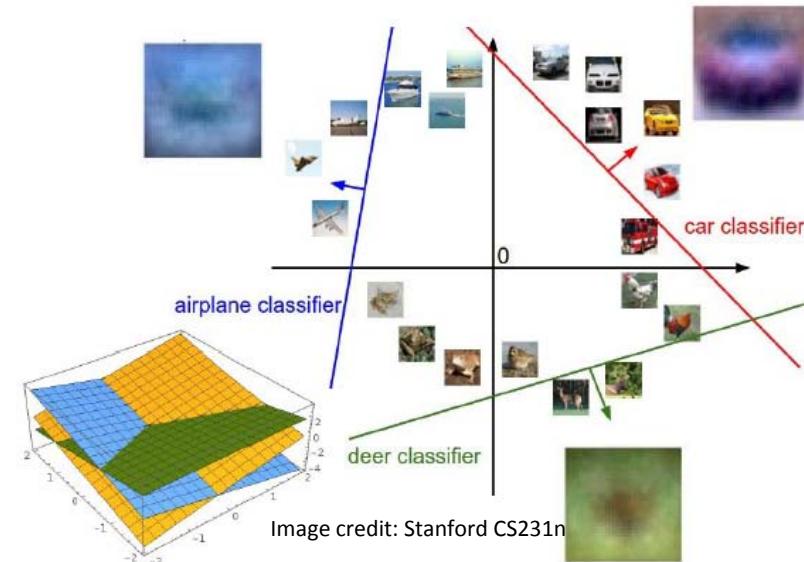
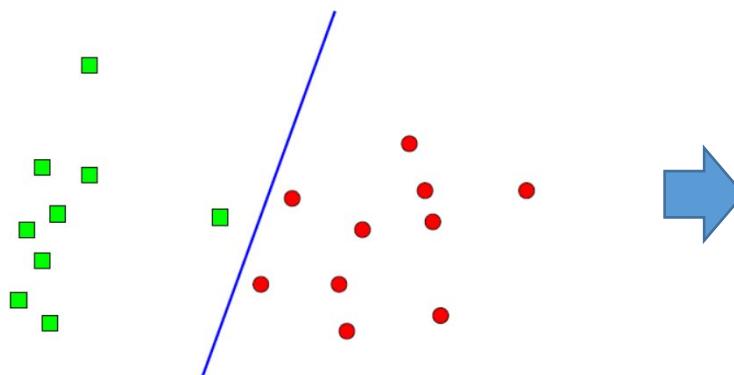
Some Remarks

- Interpreting $y = Wx + b$
 - What can we say about the learned W ?
 - The weights in W are trained by observing training data X and their ground truth Y .
 - Each column in W can be viewed as an exemplar of the corresponding class.
 - Thus, Wx basically performs **inner product** (or **correlation**) between the input x and the exemplar of each class. (Signal & Systems!)



From Binary to Multi-Class Classification

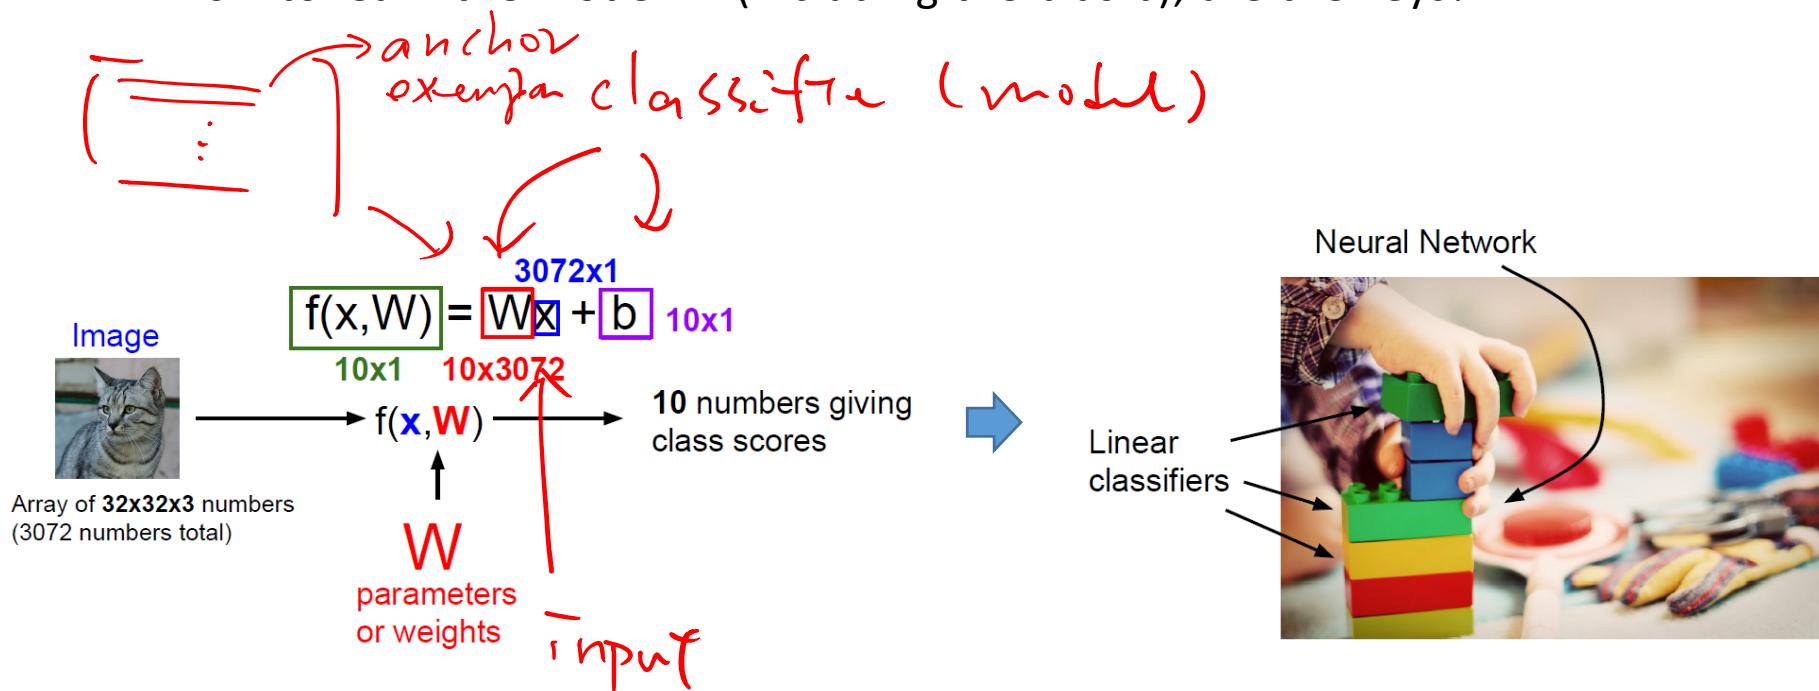
(if time permits)



- How to extend binary classifiers for solving multi-class classification?
 - 1-vs.-1 (1-against-1) vs. 1-vs.-all (1-against-rest)

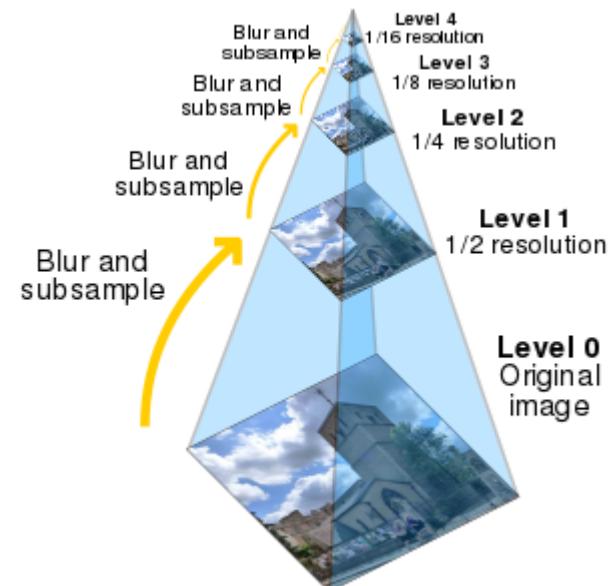
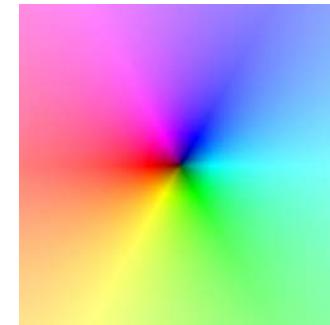
Summary for Linear Classification

- Remarks
 - Starting points for many multi-class or complex/nonlinear classifier
 - How to determine a proper loss function for matching \mathbf{y} and $\mathbf{Wx} + \mathbf{b}$, and thus how to learn the model \mathbf{W} (including the bias \mathbf{b}), are the keys!



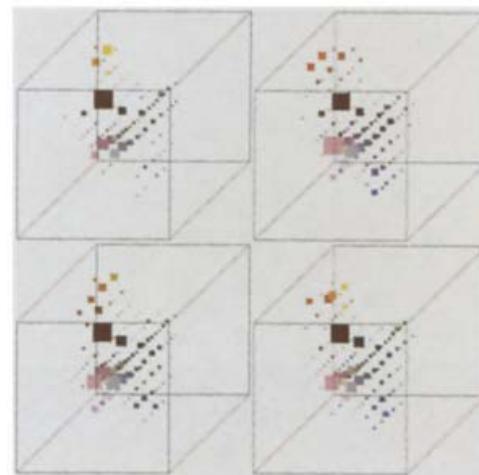
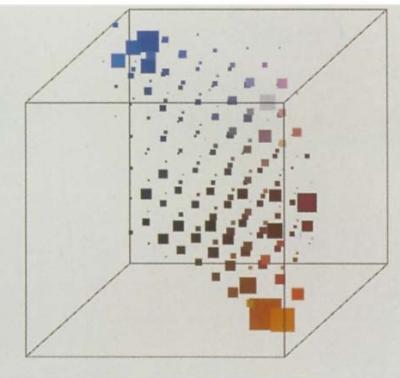
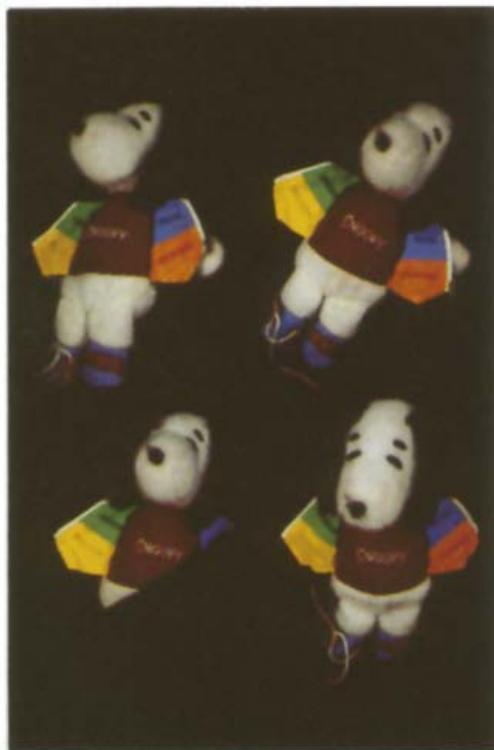
What's to Be Covered Today...

- Image Representation
 - Color Space
 - Image Filtering
 - Image Pyramid



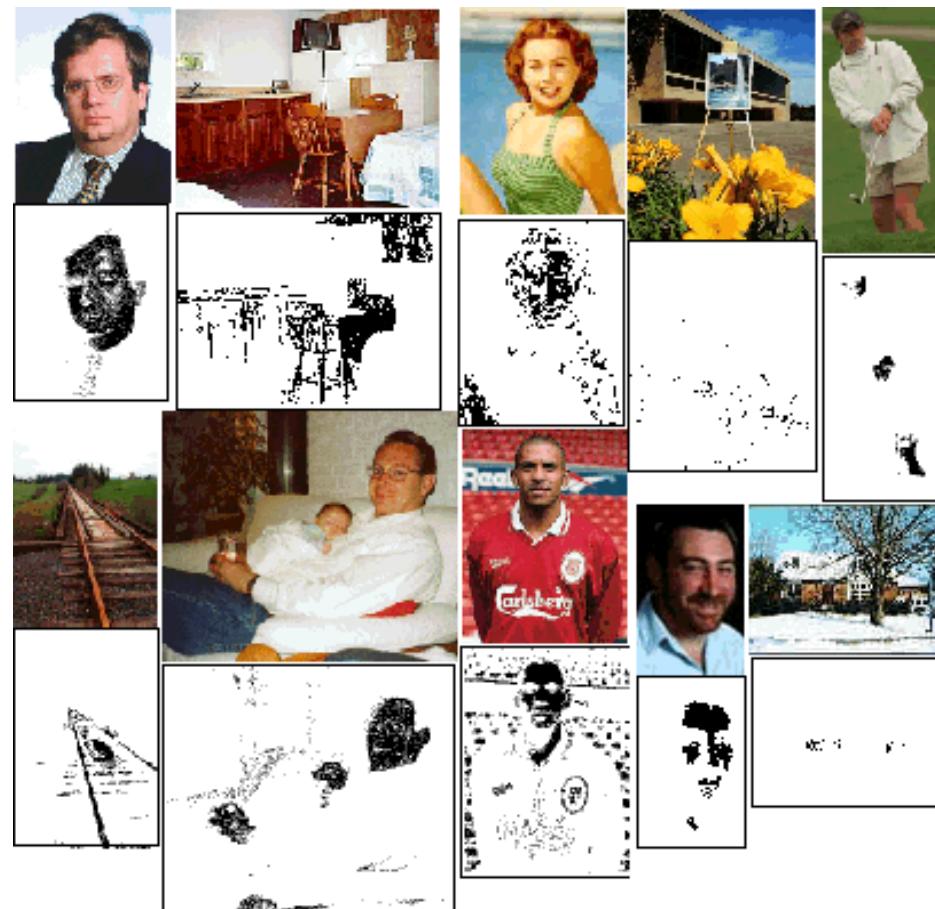
Uses of Color in Computer Vision

- Color histograms for indexing and retrieval



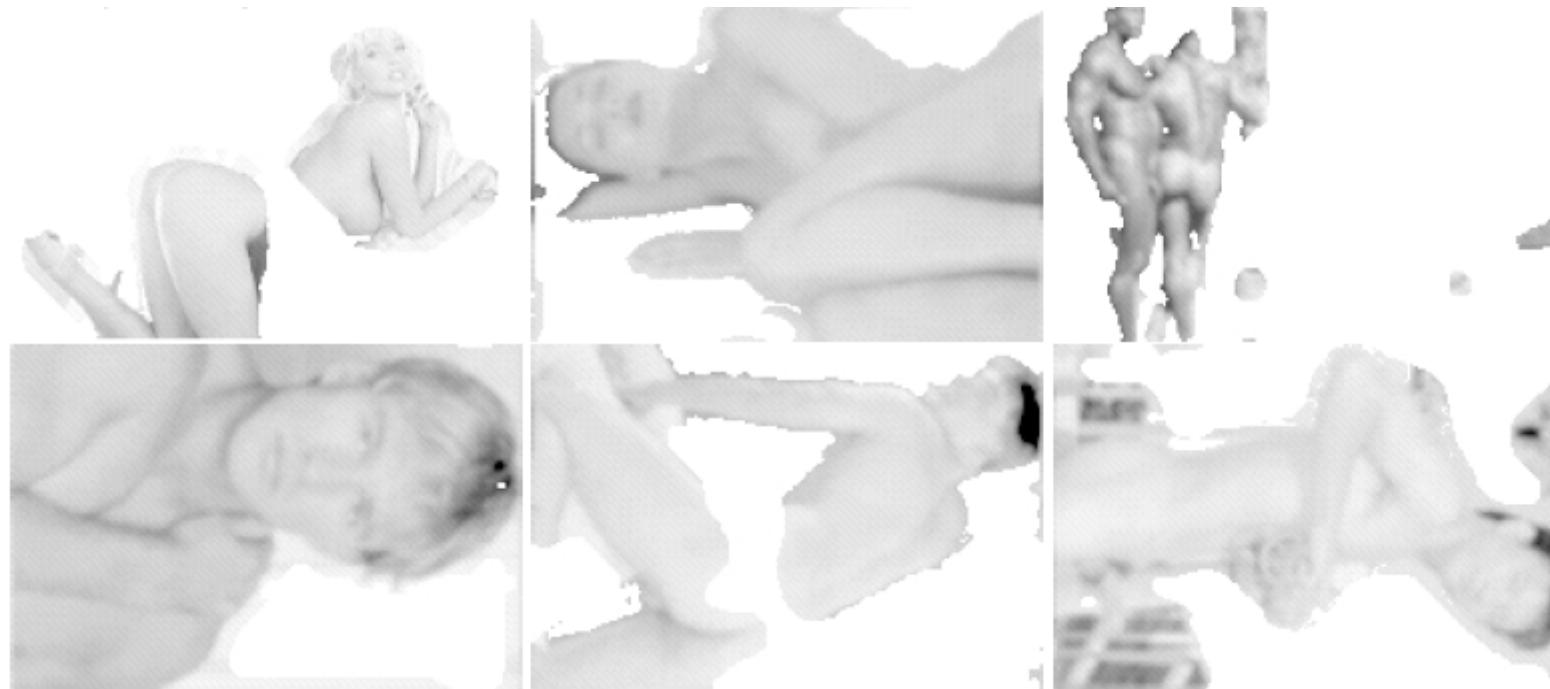
Uses of Color in Computer Vision

- Skin Detection



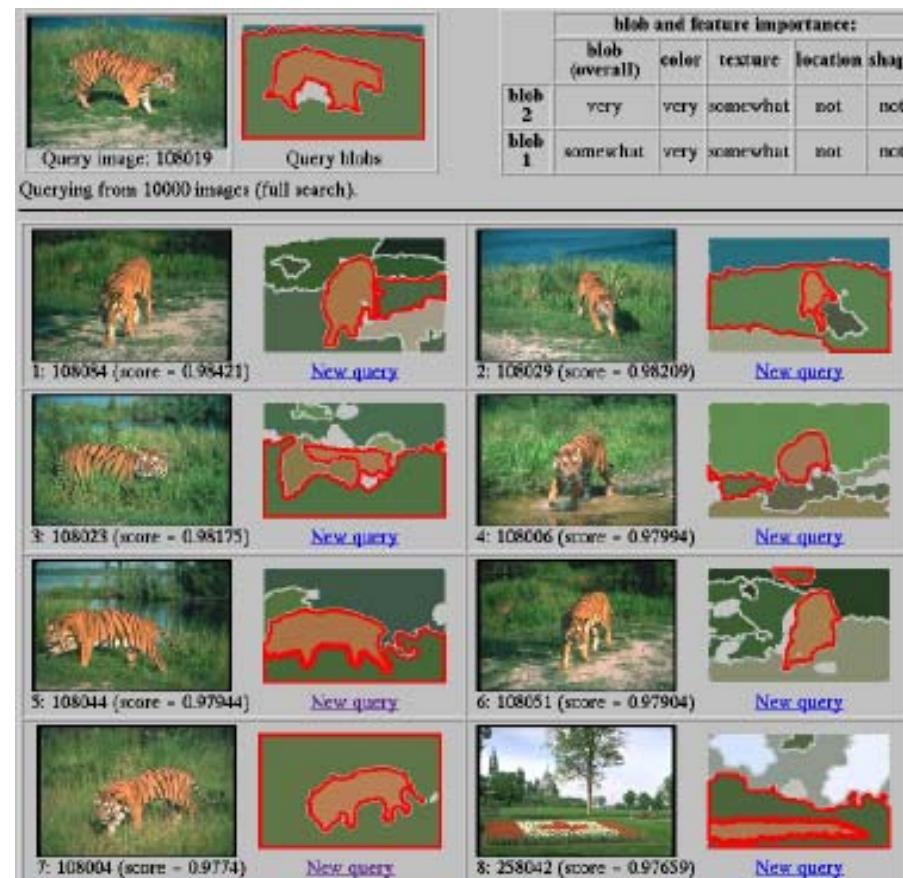
Uses of Color in Computer Vision

- Nudity detection



Uses of Color in Computer Vision

- Image Segmentation (for Detection, Retrieval, etc.)



Uses of Color in Computer Vision

- Building Appearance Models for Tracking, etc.

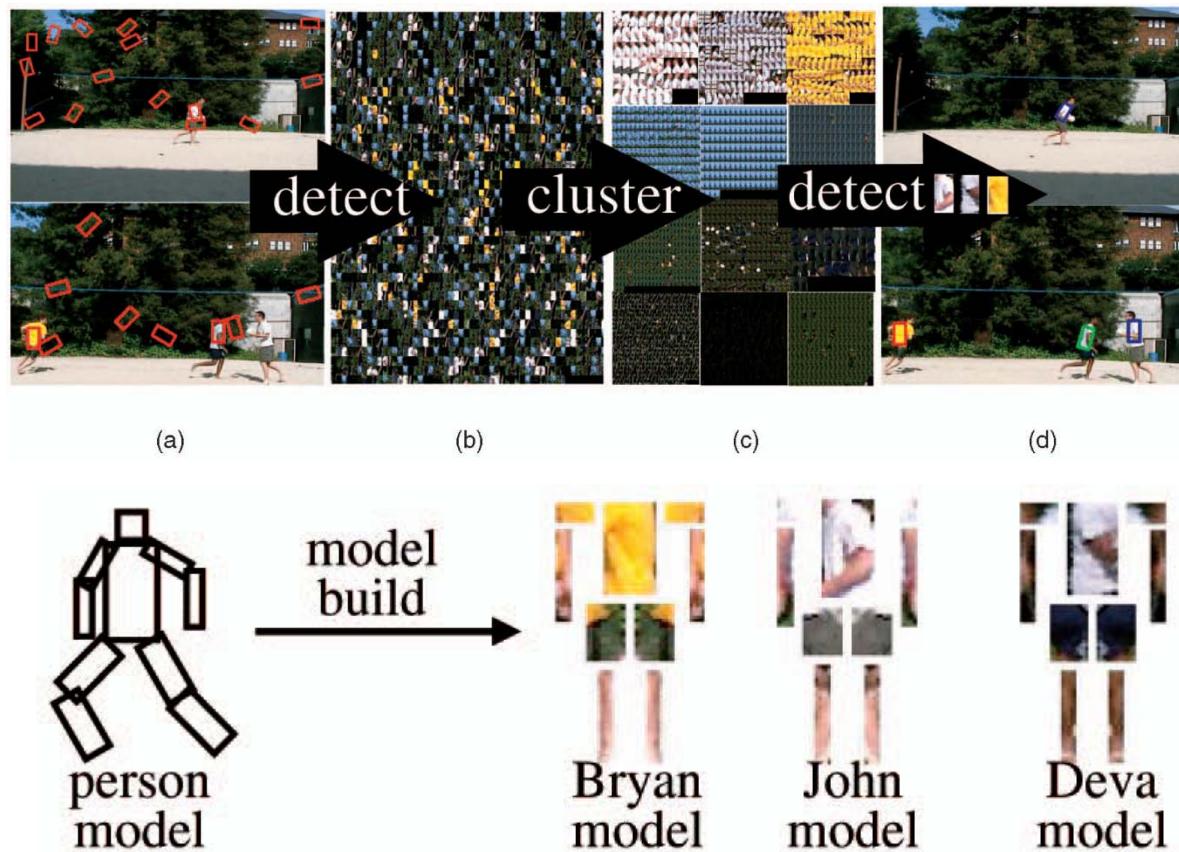


Image Representation

- Note that, since we focus on learning-based vision, we will not cover geometry, projection, camera models, etc. in this course.

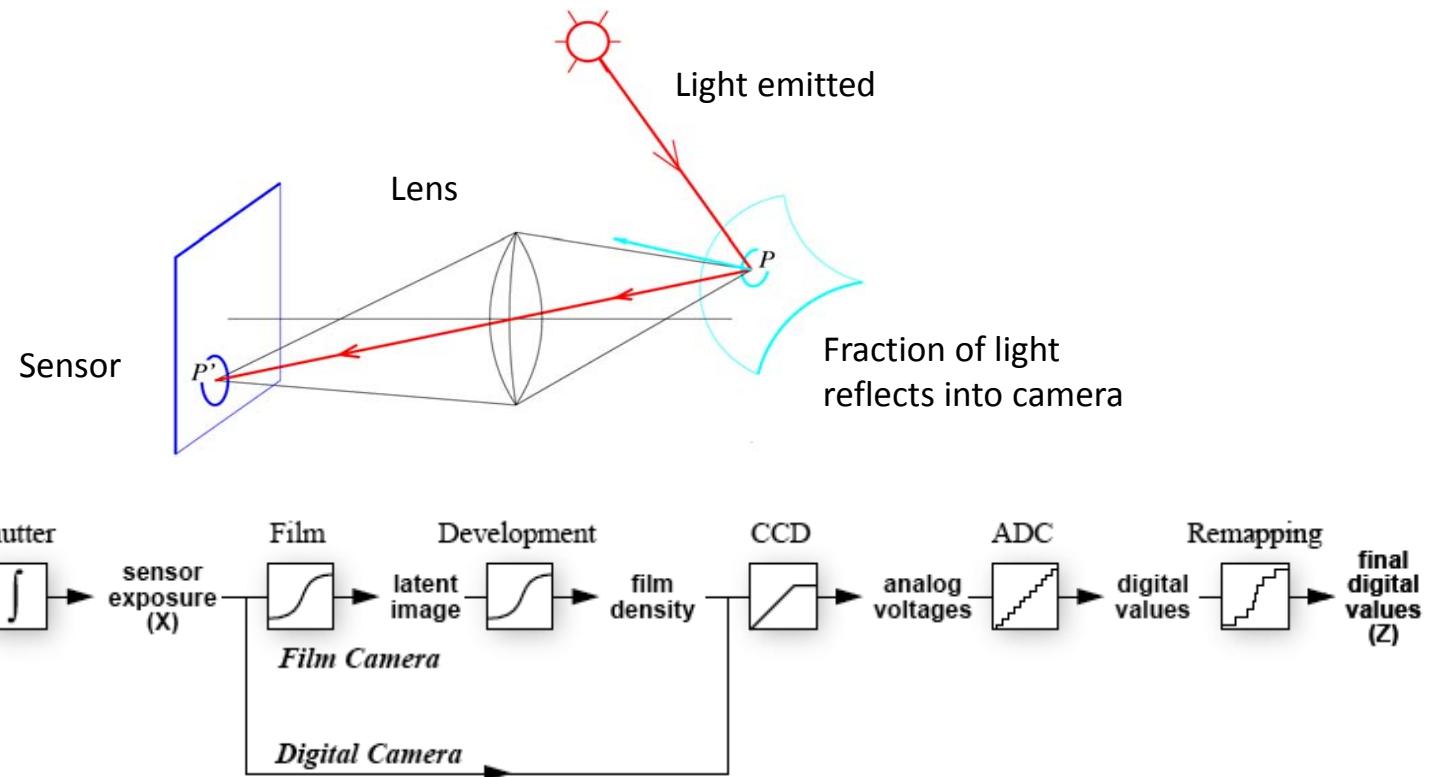
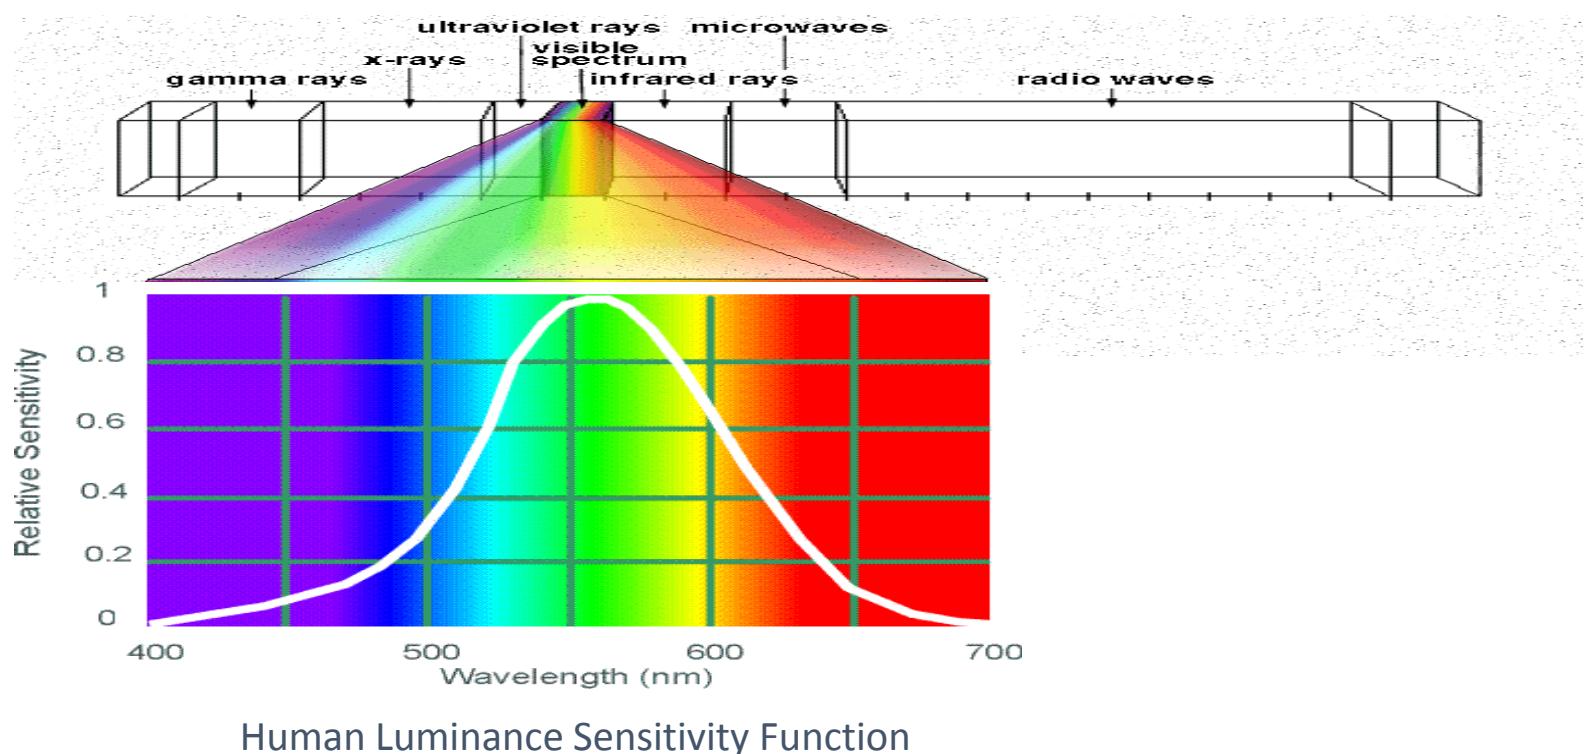


Image credit: Rob Fergus

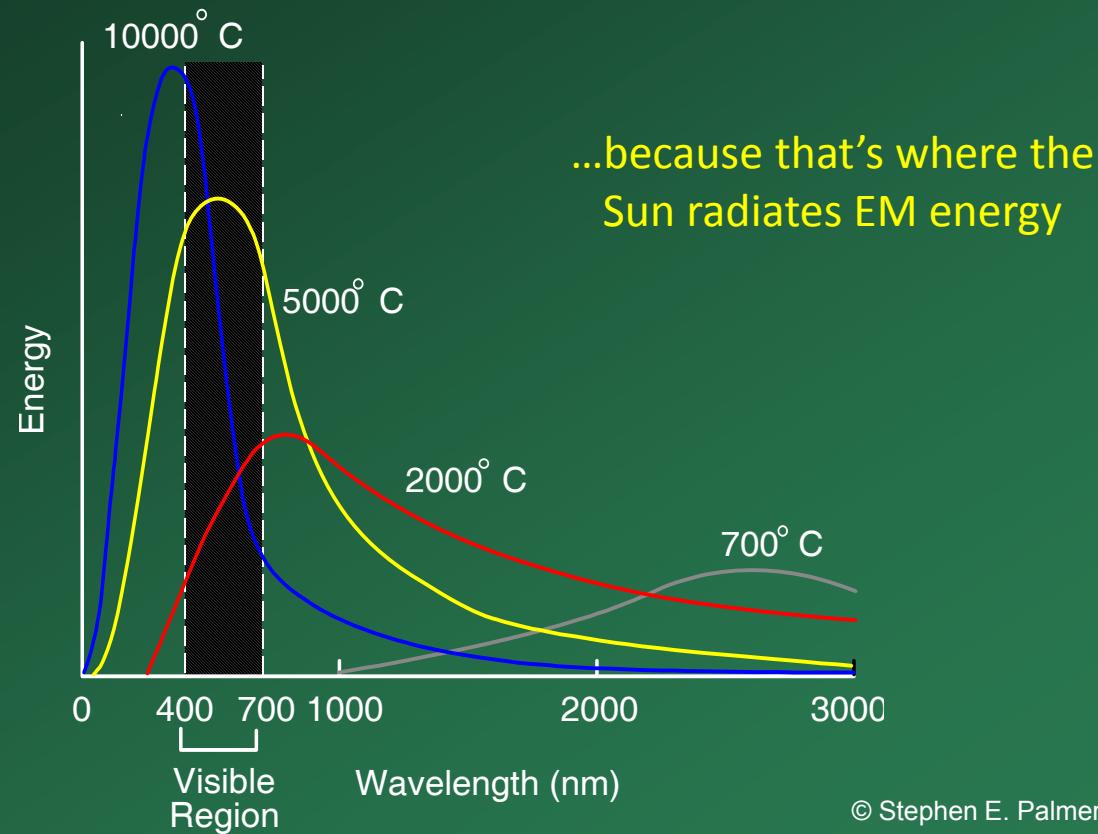
The Physics of Light

- Light
 - Electromagnetic energy whose wavelength is between 400 nm and 700 nm. ($1 \text{ nm} = 10^{-9} \text{ meter}$)



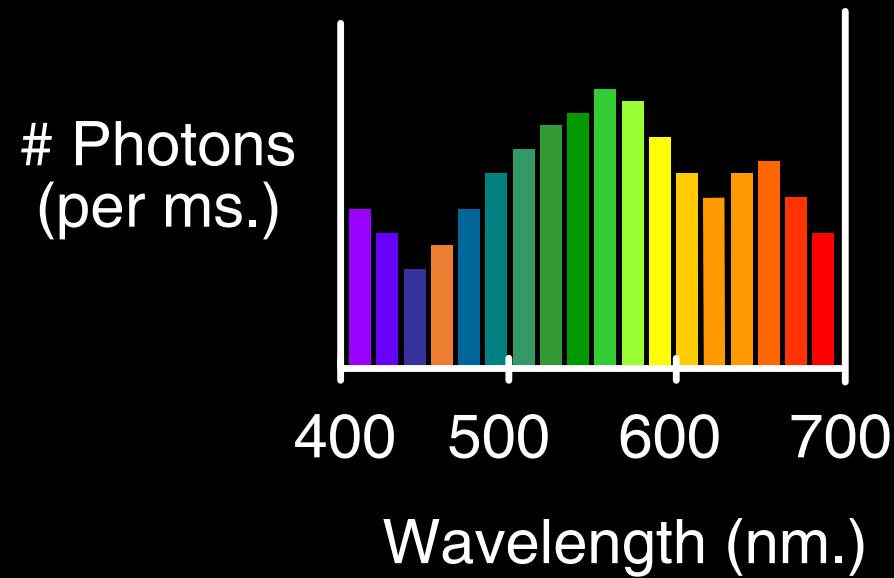
Visible Light

Why do we see light of these wavelengths?



The Physics of Light

Any patch of light can be completely described physically by its spectrum: the number of photons (per time unit) at each wavelength 400 - 700 nm.

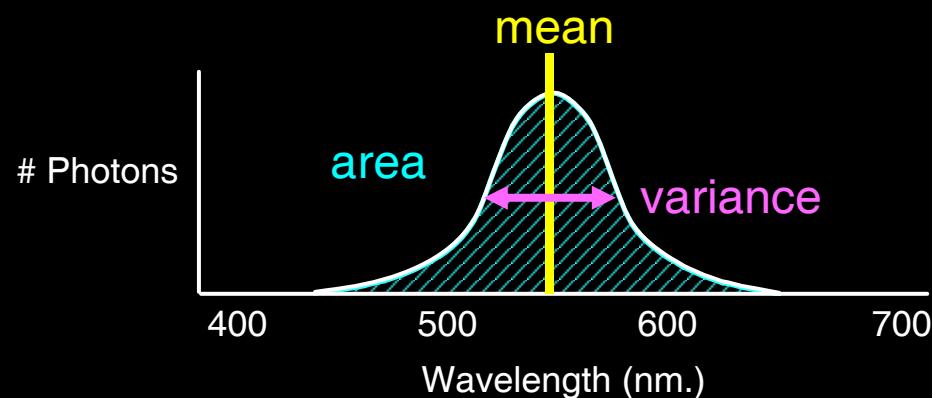


The Psychophysical Correspondence

There is no simple functional description for the perceived color of all lights under all viewing conditions, but

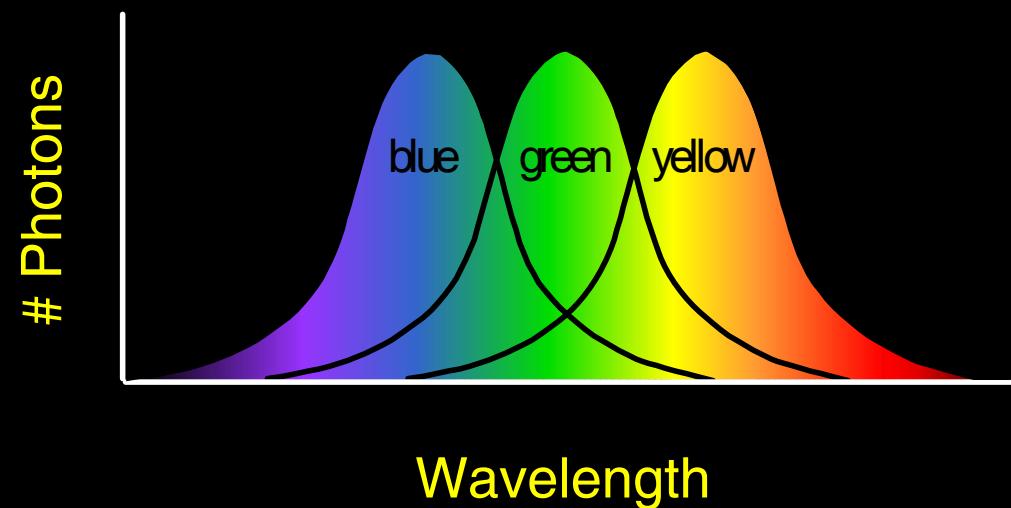
A helpful constraint:

Consider only physical spectra with normal distributions



The Psychophysical Correspondence

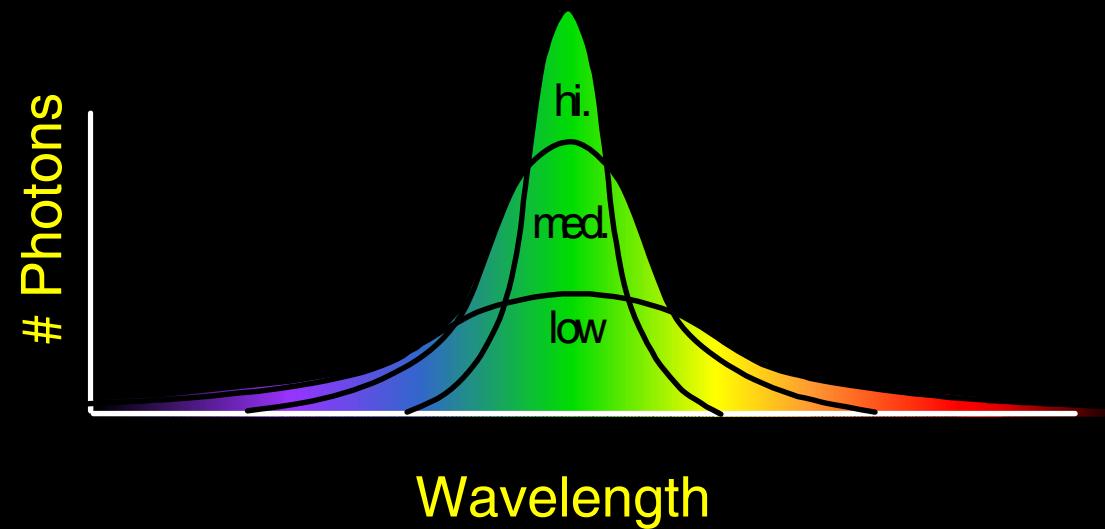
Mean \longleftrightarrow Hue



© Stephen E. Palmer, 2002

The Psychophysical Correspondence

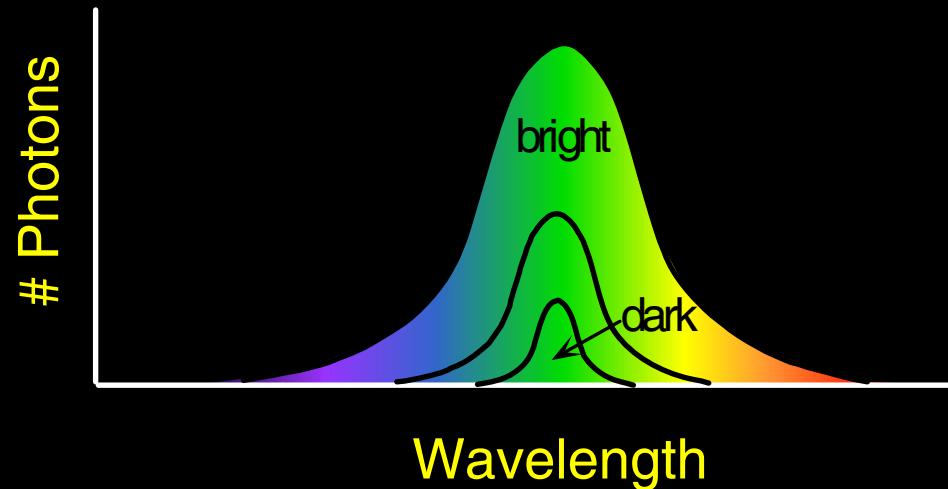
Variance \longleftrightarrow Saturation



© Stephen E. Palmer, 2002

The Psychophysical Correspondence

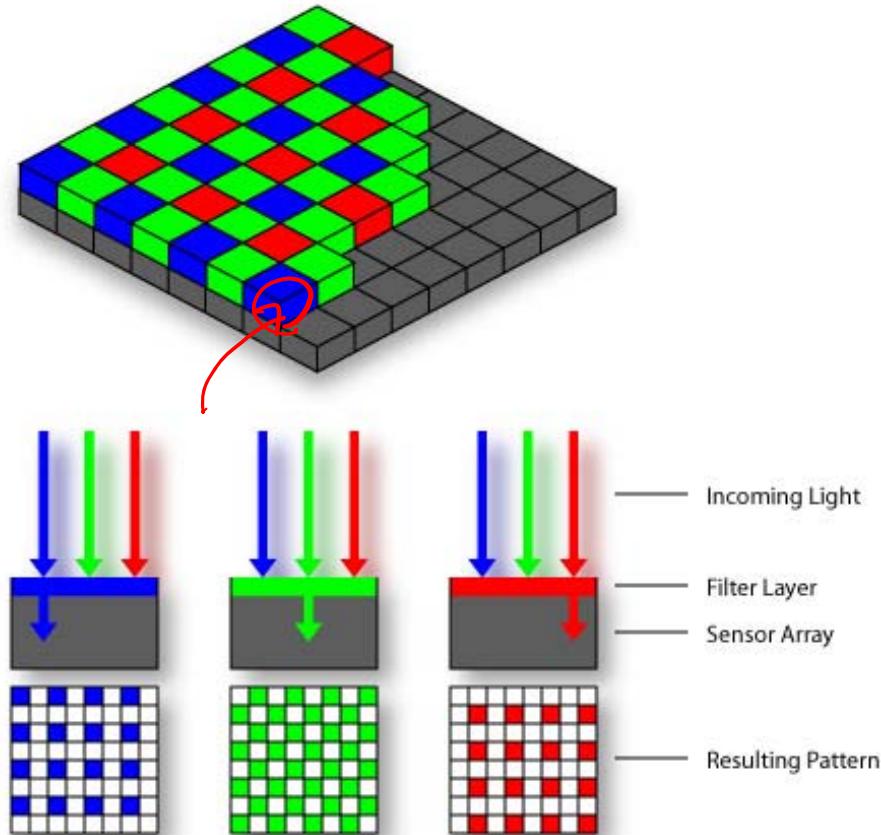
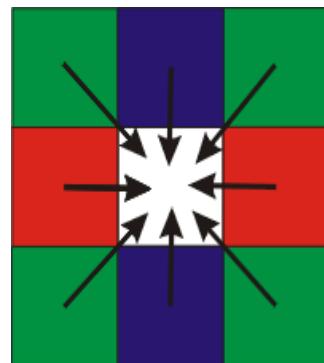
Area \longleftrightarrow Brightness



© Stephen E. Palmer, 2002

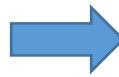
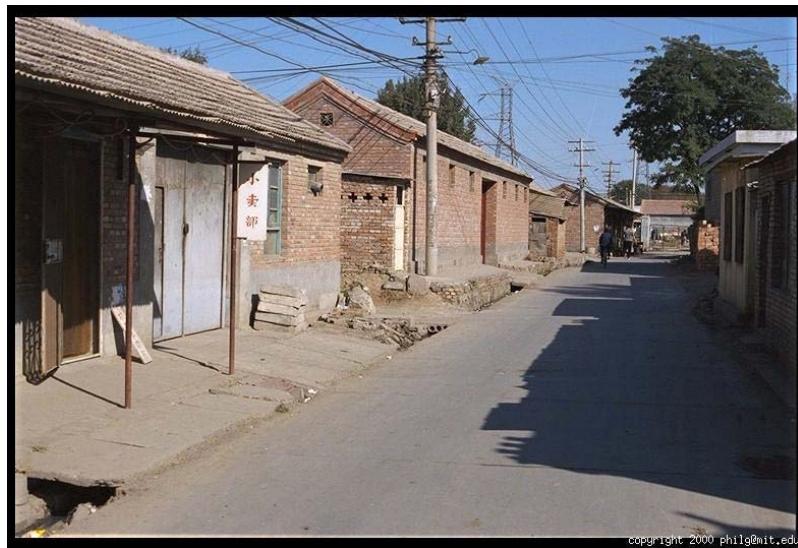
Practical Color Sensing

- Bayer Grid/Pattern



- Estimate RGB at 'G' cells from neighboring values

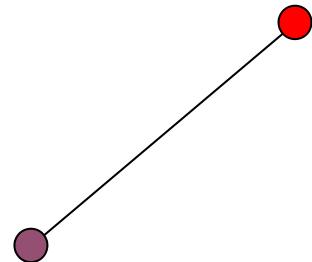
Example Color Image



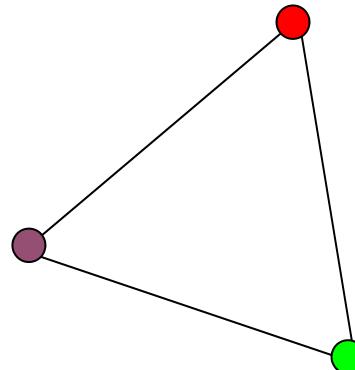
Color Space



- How can we represent color?
- Linear color spaces
 - Defined by a choice of three primaries
 - The coordinates of a color are given by the weights of each primary.



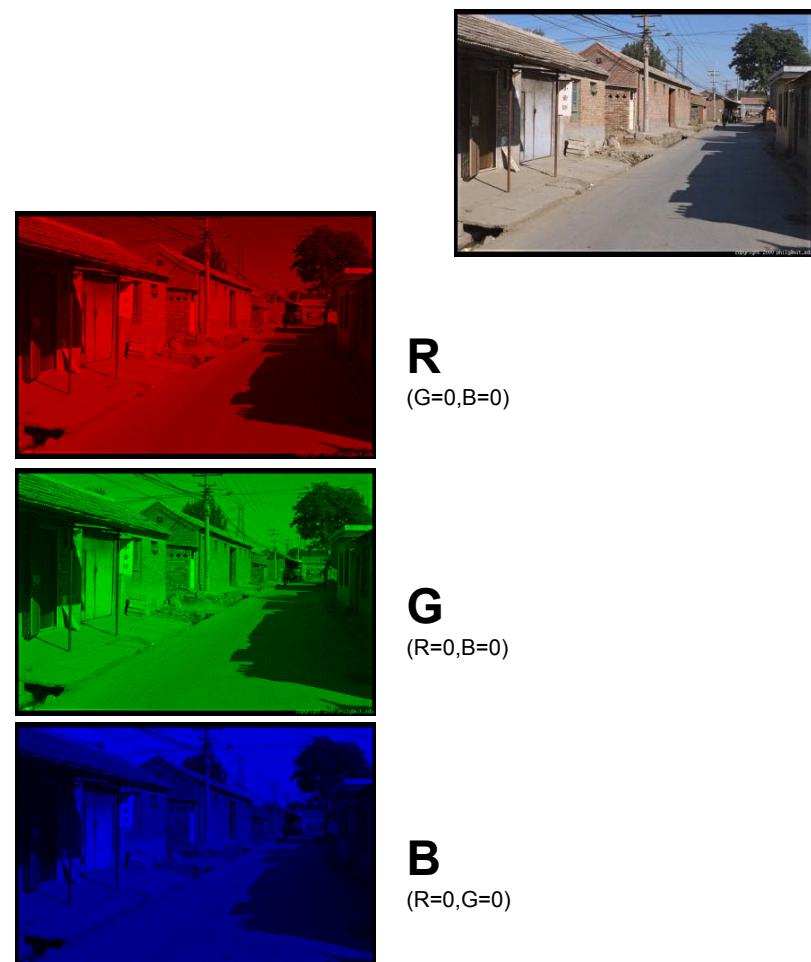
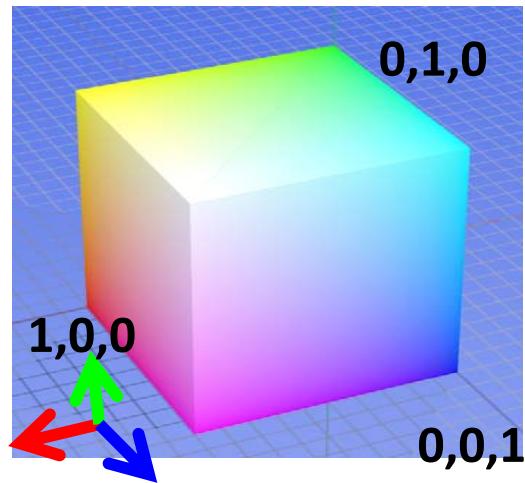
Mixing two lights produces colors that lie along a straight line in color space.



Mixing three lights produces colors that lie within the triangle defined in the color space.

Example Color Space: RGB

- Default Color Space

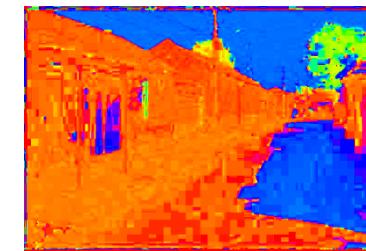
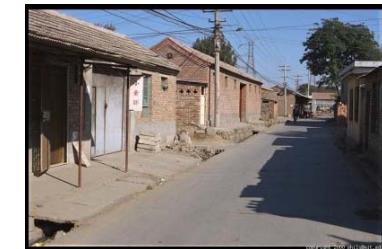
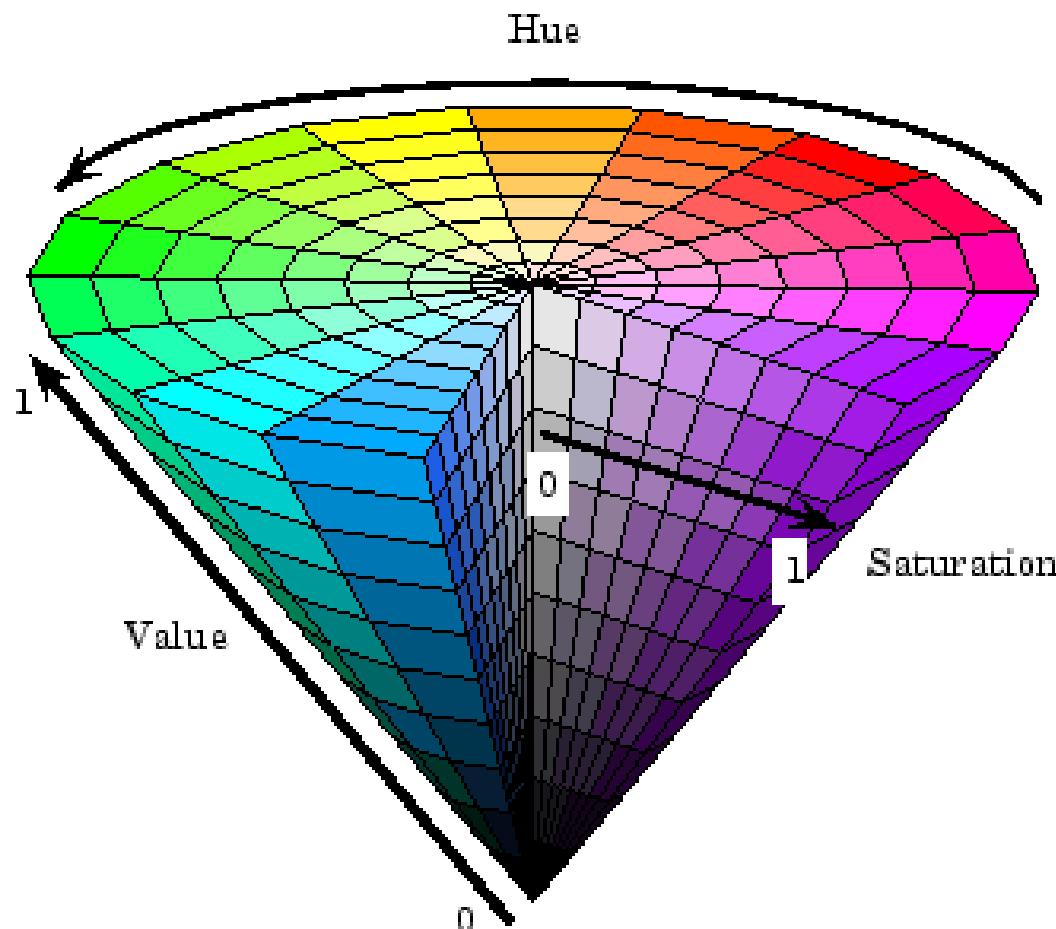


- Remarks
 - Easy for devices
 - But not perceptual
 - Where do the grays live?
 - Where is hue and saturation?

Image from: http://en.wikipedia.org/wiki/File:RGB_color_solid_cube.png

Example Color Space: HSV

- Intuitive Color Space



H
 $(S=1, V=1)$



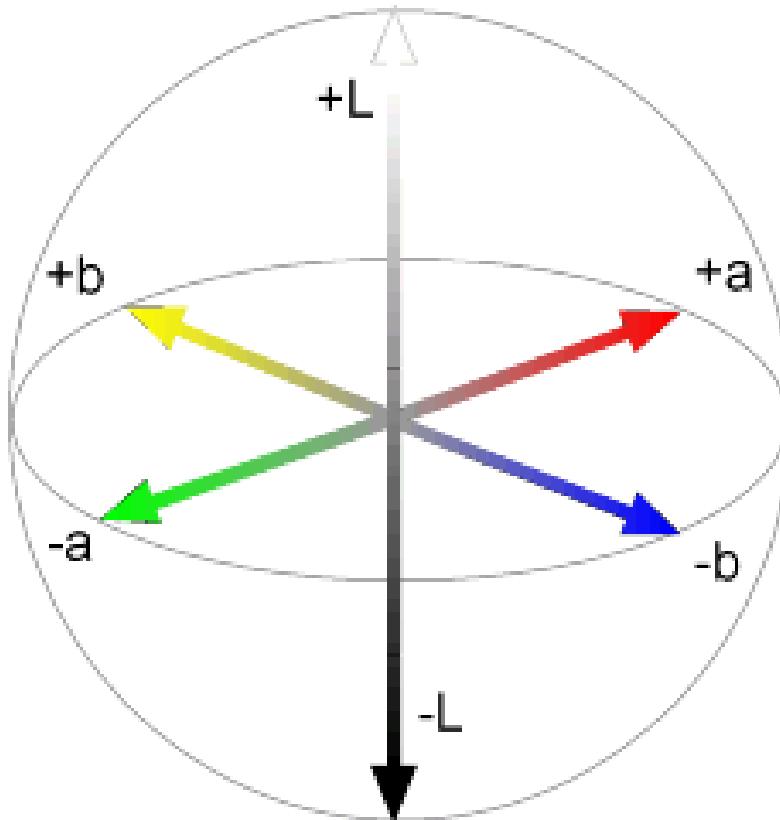
S
 $(H=1, V=1)$



V
 $(H=1, S=0)$

Example Color Space: L*a*b*

- “Perceptually uniform/linear” color space



L
($a=0, b=0$)



a
($L=65, b=0$)

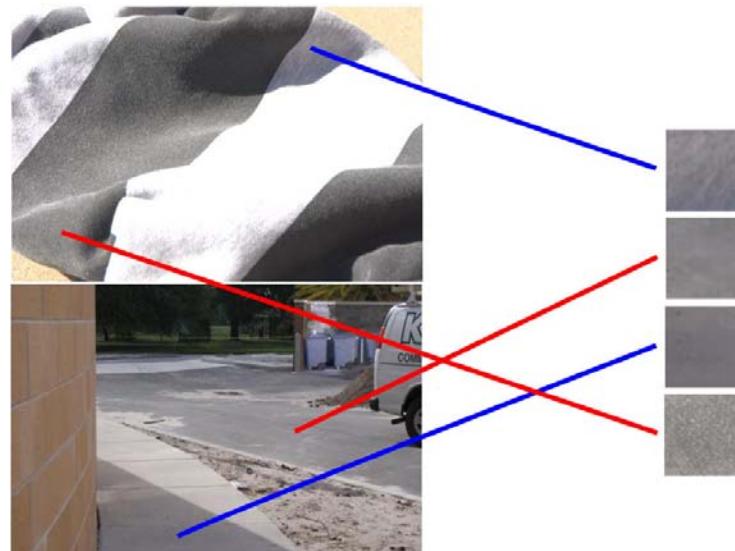


b
($L=65, a=0$)



Light

- A pixel's brightness is determined by
 - Light source (strength, direction, color)
 - Surface orientation
 - Surface material and albedo
 - Reflected light and shadows from surrounding surfaces
 - Gain on the sensor
- A pixel's brightness tell us nothing but itself.



Slide credit: J.-B. Huang

Light (cont'd)

- Yet, we are able to interpret images from lightness.
 - For nearby pixels/points, most factors do not change much.
 - The information is mainly contained in *local differences* of brightness



copyright 2000 philg@mit.edu

What is this?





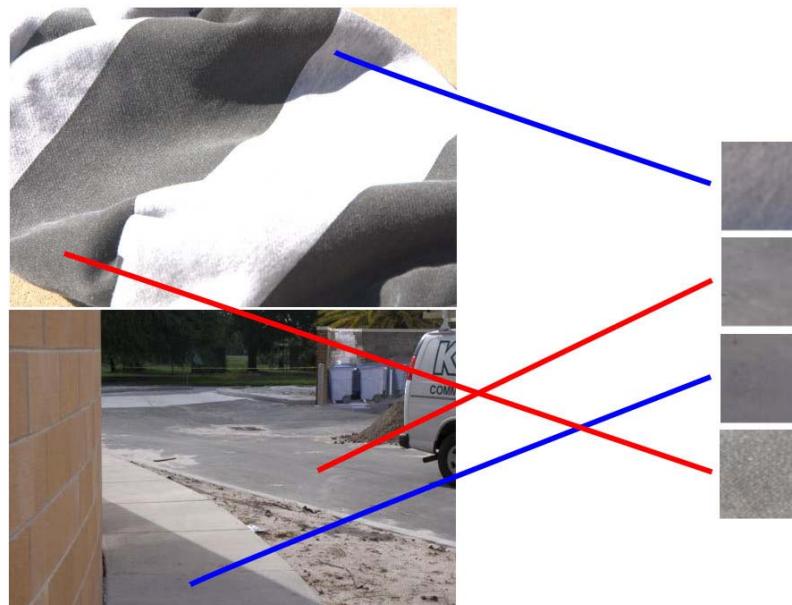
What differences in intensity tell us about shape?

- Changes in surface normal
- Texture
- Proximity
- Indents and bumps
- Grooves and creases



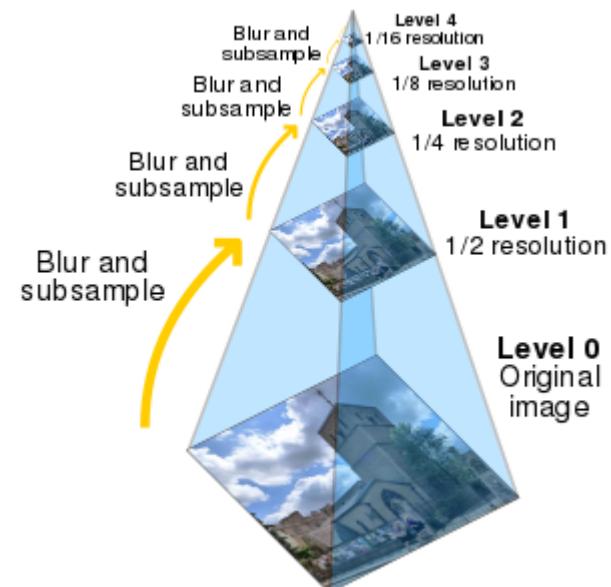
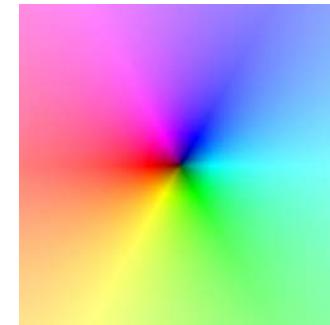
Color Constancy

- Interpret surface in terms of albedo or “true color”, rather than observed intensity.
 - Humans are good at it.
 - Computers are not nearly as good...

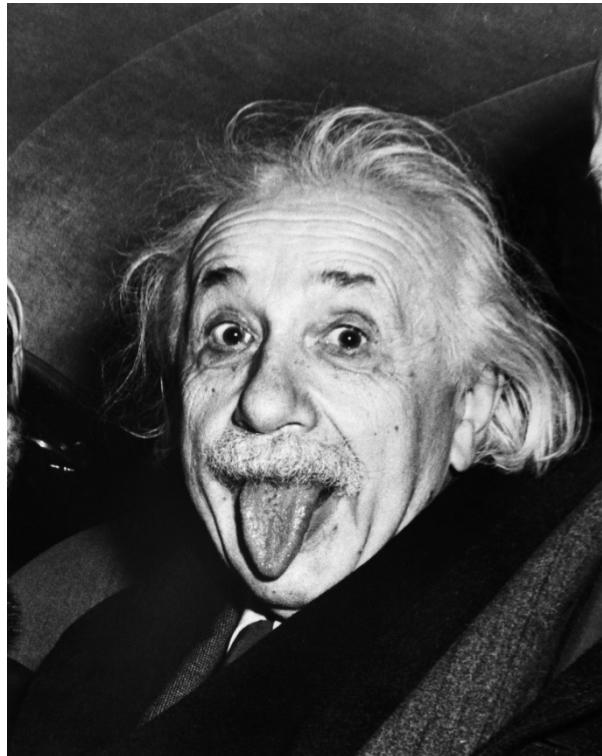


What's to Be Covered Today...

- Image Representation
 - Color Space
 - Image Filtering
 - Image Pyramid



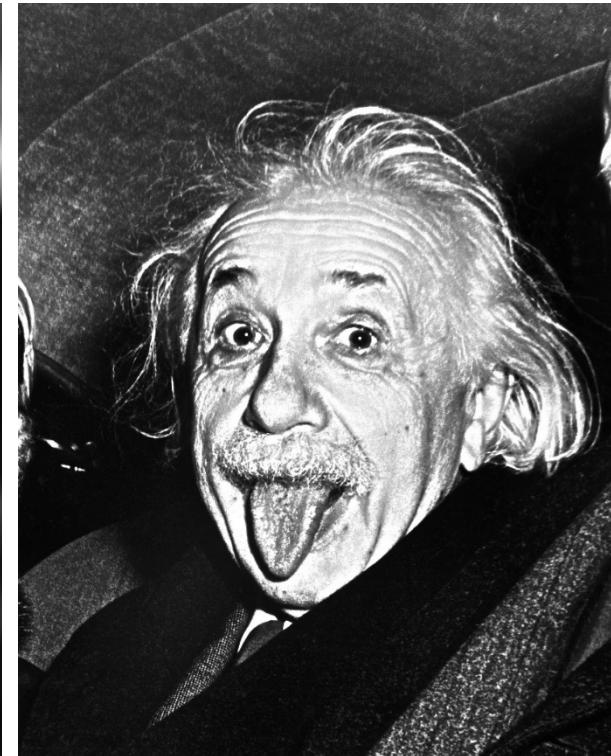
Why Image Filtering?



Input



Smoothing



Sharpening

Why Image Filtering?

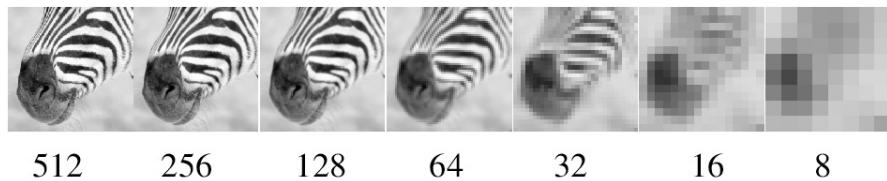


Image Pyramid

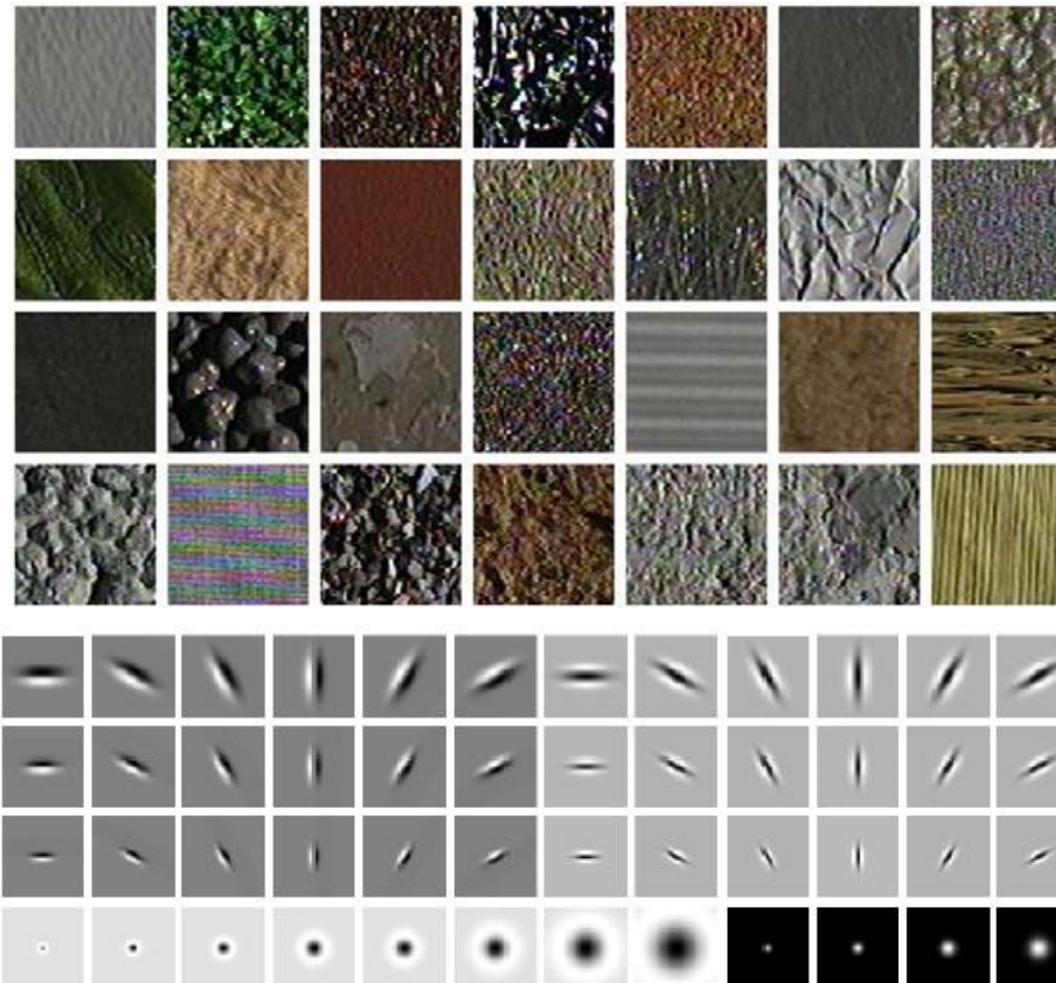


Image interpolation/resampling



Why Image Filtering?

Representing textures with filter banks



LM filter bank. Code [here](#)

What's Image Filtering?

- Image Filtering
 - For each pixel of interest, compute the function of local neighborhood and output the new value.
 - Generally (while not always true), same function applied to each position
 - Typically (while not required), output and input images are of the same size.

10	5	3
4	5	1
1	1	7

Local image data

Some function



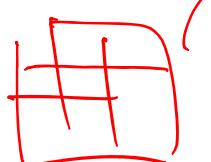
	7	

Modified image data

What's Image Filtering?

- Linear filtering

- function is a weighted sum/difference of pixel values



10	5	3
4	6	1
1	1	8

Local image data

0	0	0
0	0.5	0
0	1	0.5

Filter/kernel

		8

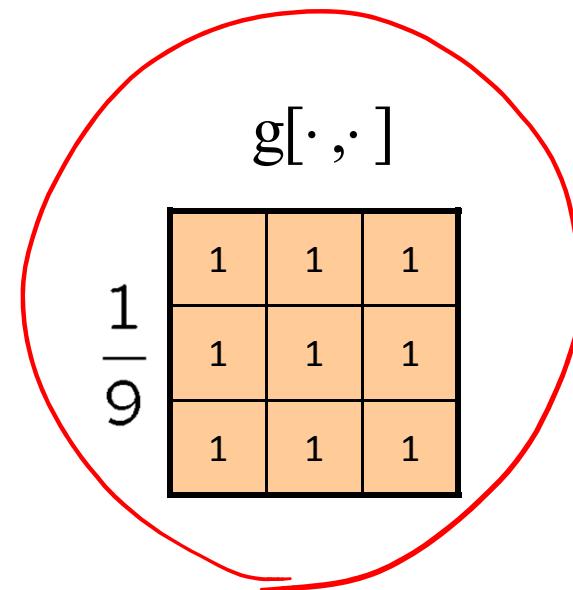
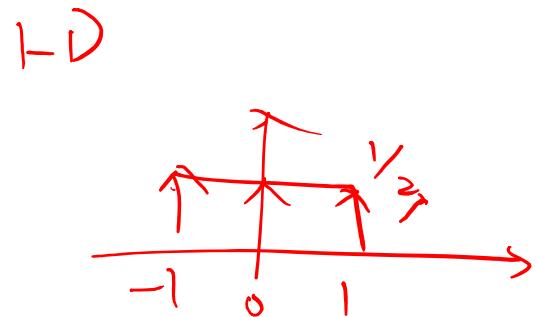
Modified image data

- Example applications

- Enhance images
 - Denoise, smooth, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching



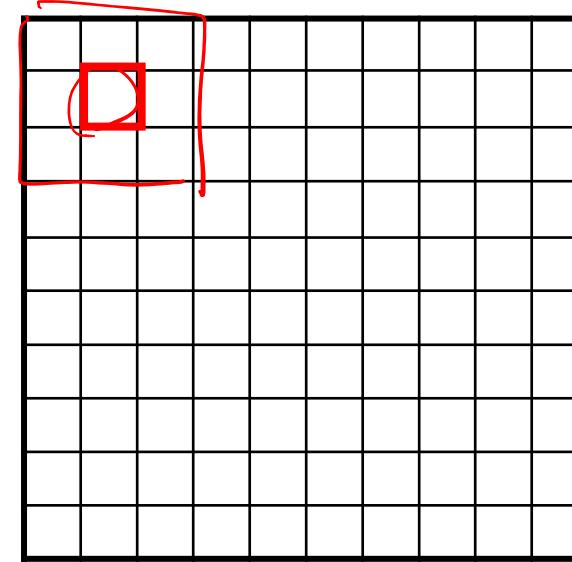
Example: Image Filtering with a Box Filter



img
 $f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g[\cdot, \cdot] \frac{1}{9}$
 $h[.,.]$

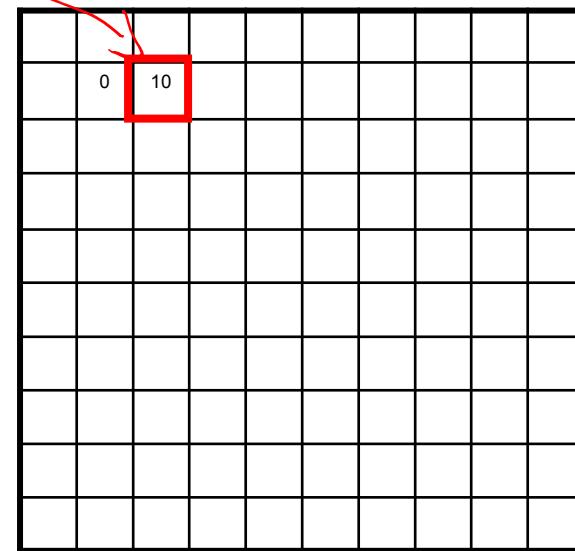
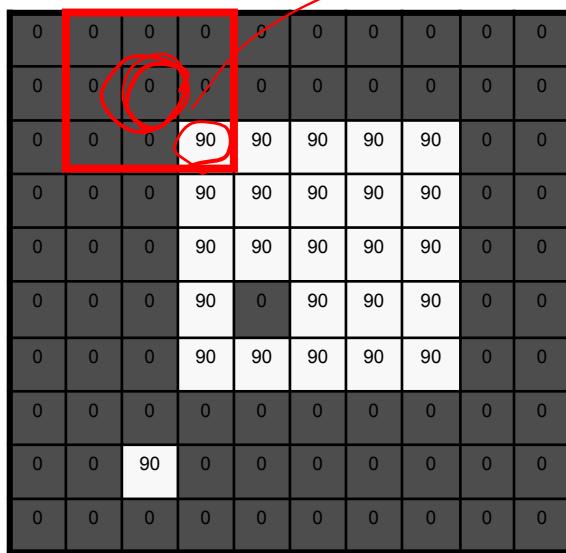


$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9}$$

$f[\cdot, \cdot]$

$h[\cdot, \cdot]$



$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10	20									

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

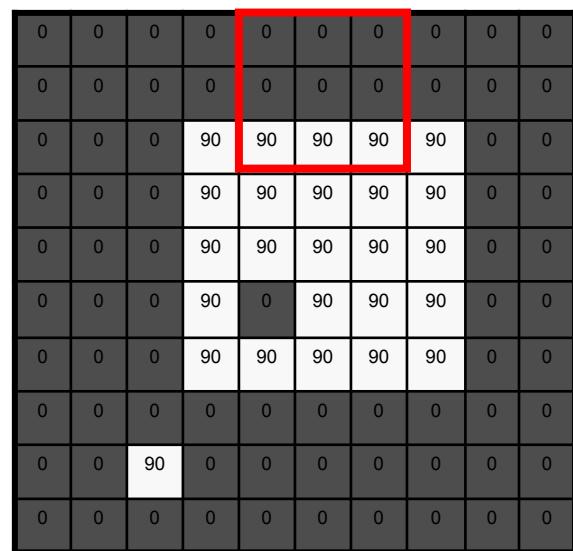
$h[.,.]$

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

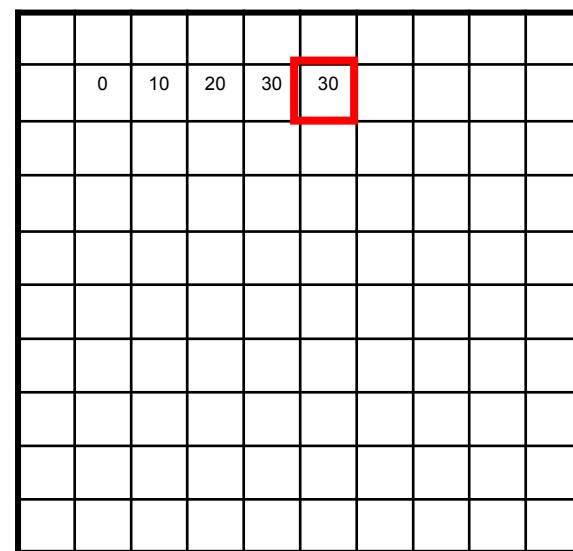
$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$



$$h[.,.]$$



$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10	20	30	30							

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

$$h[.,.]$$

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

$$f[\cdot, \cdot]$$

$$g[\cdot, \cdot]$$

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$h[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

0	10	20	30	30	30	20	10	
0	20	40	60	60	60	40	20	
0	30	60	90	90	90	60	30	
0	30	50	80	80	90	60	30	
0	30	50	80	80	90	60	30	
0	20	30	50	50	60	40	20	
10	20	30	30	30	30	20	10	
10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

$k, l \in \{-1, 0, 1\}$

Example: Image Filtering with a Box Filter

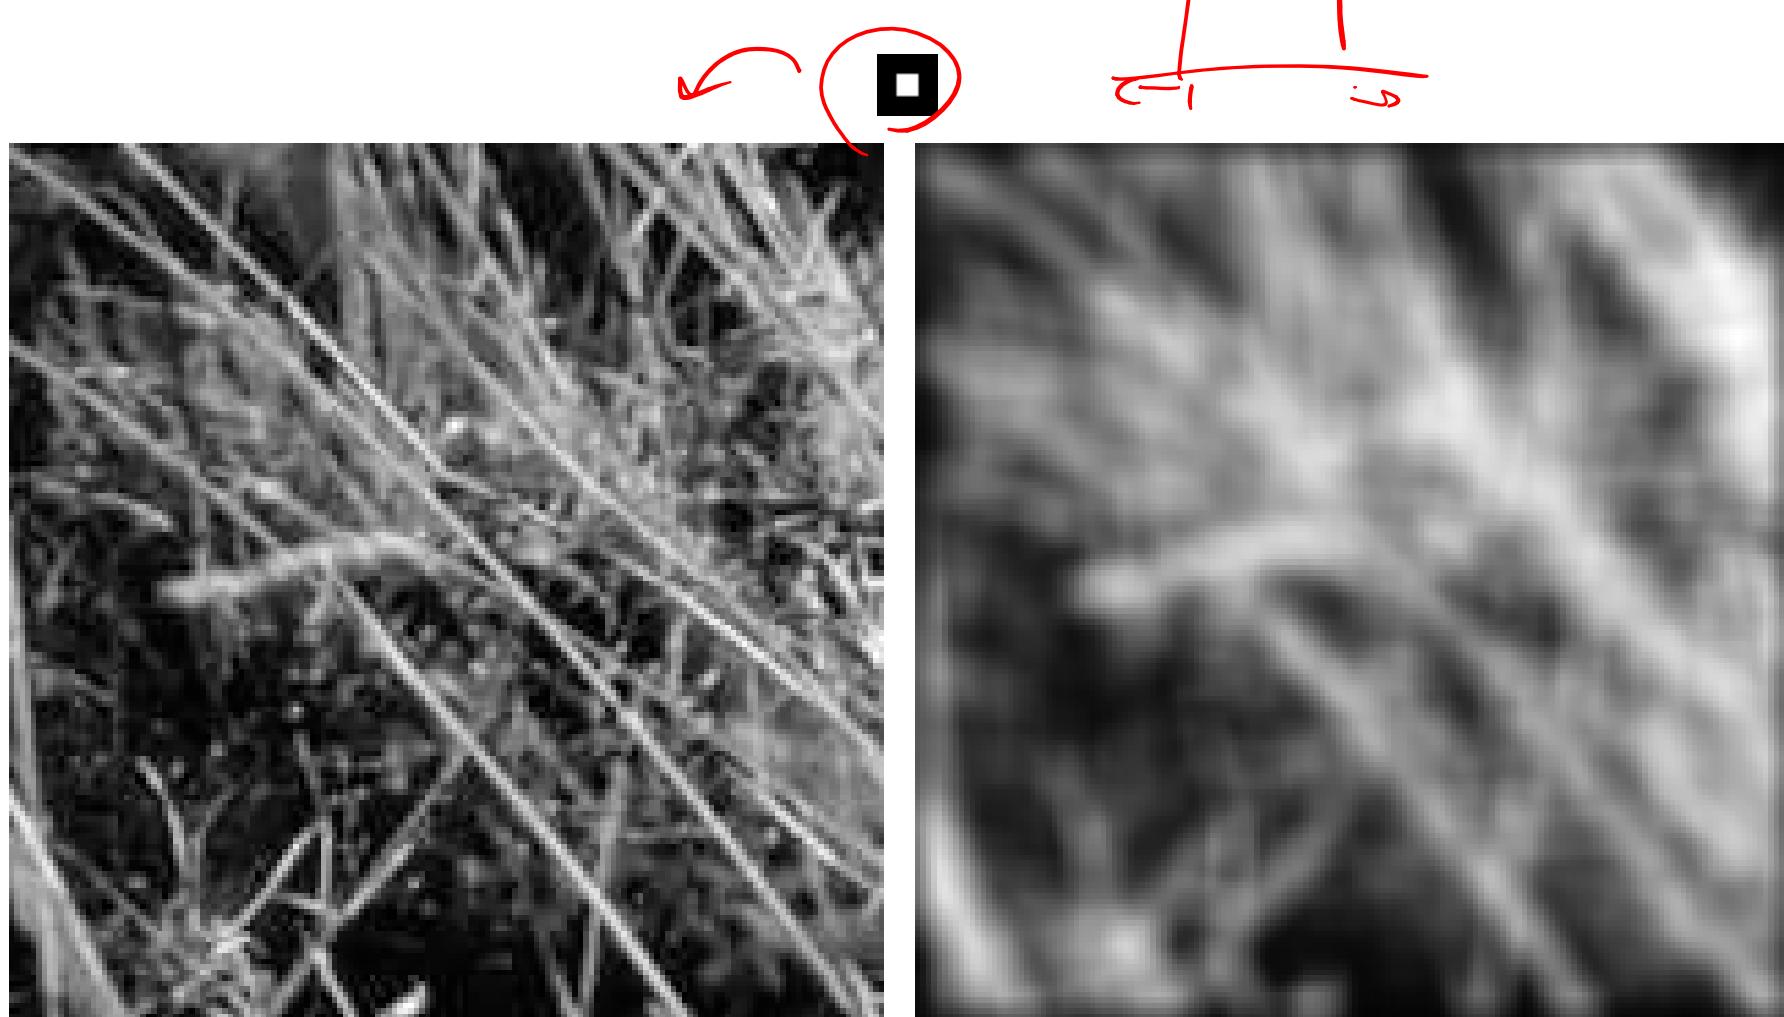
- What does Box Filter do?
 - Replaces each pixel with an average of its neighborhood
 - Performs image smoothing
 - Achieves image denoising?



$g[\cdot, \cdot]$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

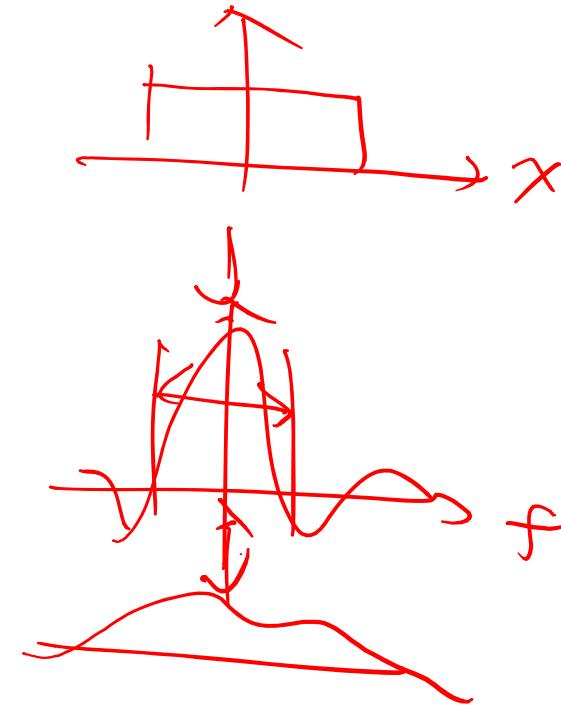
Smoothing with a Box Filter



Properties of Smoothing Filters

- Smoothing
 - Values positive
 - Sum to 1 → constant regions same as input
 - Amount of smoothing proportional to the mask/filter size
 - Remove “high-frequency” components, i.e., **low-pass filter (LPF)**

$$g[\cdot, \cdot]$$
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



Slide credit: K. Grauman

Correlation Filtering

Say the averaging window size is $2k+1 \times 2k+1$:

$$G[i, j] = \frac{1}{(2k+1)^2} \sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]$$

img

Attribute uniform weight to each pixel

neighbors

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v]$$

Correlation Filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted

$$G = H \otimes F$$

Note that the filter “kernel” or “mask” $H[u, v]$ is the prescription for the pixel weights for linear combination.

Filtering an Impulse Signal

What is the result of filtering the impulse signal (image) F with an arbitrary kernel H ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	1	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$$F[x, y]$$



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

				a		
				?		

$$G[x, y]$$

Convolution

- Convolution:

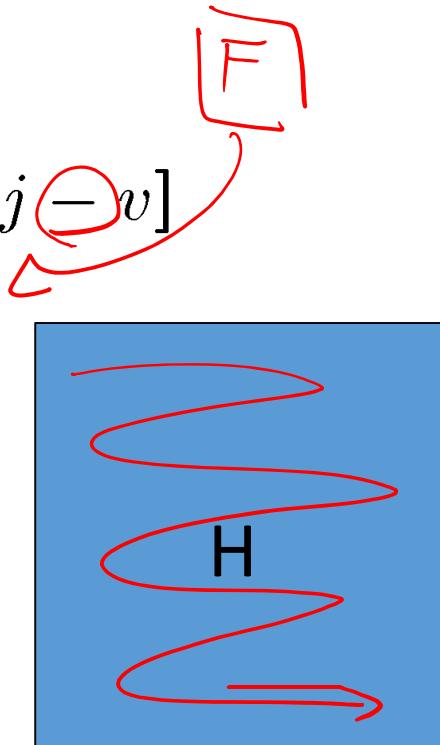
- Flip the filter in both dimensions (bottom to top, right to left)
- Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

↑

Notation for convolution operator



Convolution vs. Correlation

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v] \quad G = \text{conv2}(H, F);$$

$$G = H \star F$$

Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v] \quad G = \text{filter2}(H, F); \text{ or } G = \text{imfilter}(F, H);$$

$$\uparrow$$

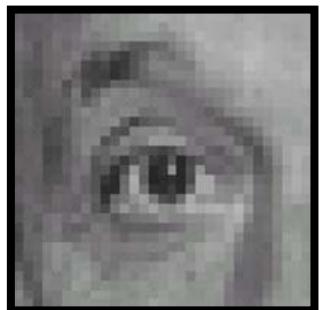
$$G = H \otimes F$$

$$\uparrow$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

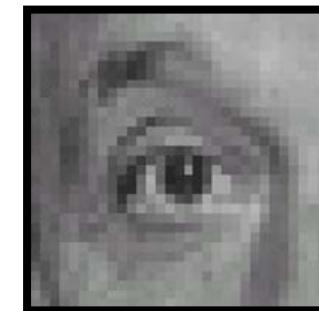
Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

f



Filtered
(no change)

↙ ↗

A red hand-drawn arrow indicating the flow from the original image to the filtered image.

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left or right?
By 1 pixel

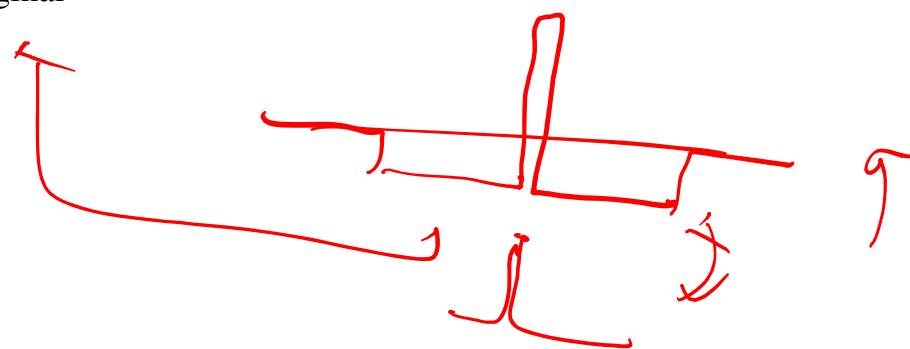
Practice with linear filters


$$\text{Original}$$

$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix} - \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$

(Note that filter sums to 1)

?



Slide credit: David Lowe (UBC)

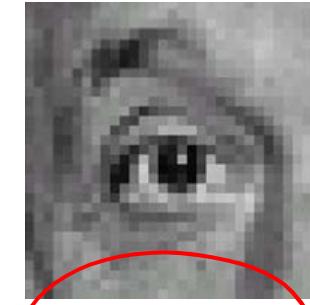
Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

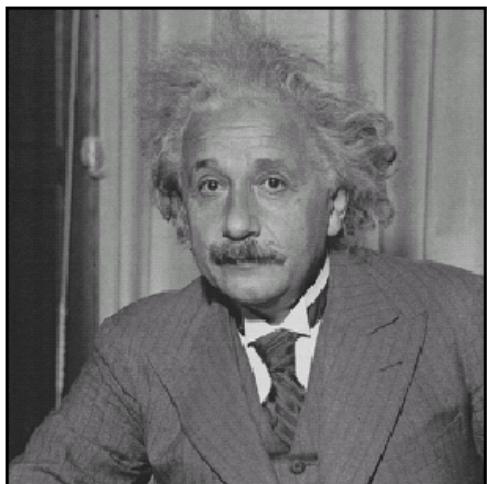
$$- \frac{1}{9} \begin{array}{|ccc|} \hline 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ \hline \end{array}$$



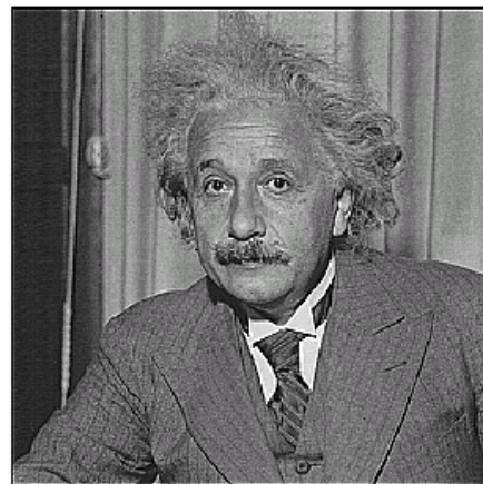
Sharpening filter

Accentuates differences with local average

Sharpening

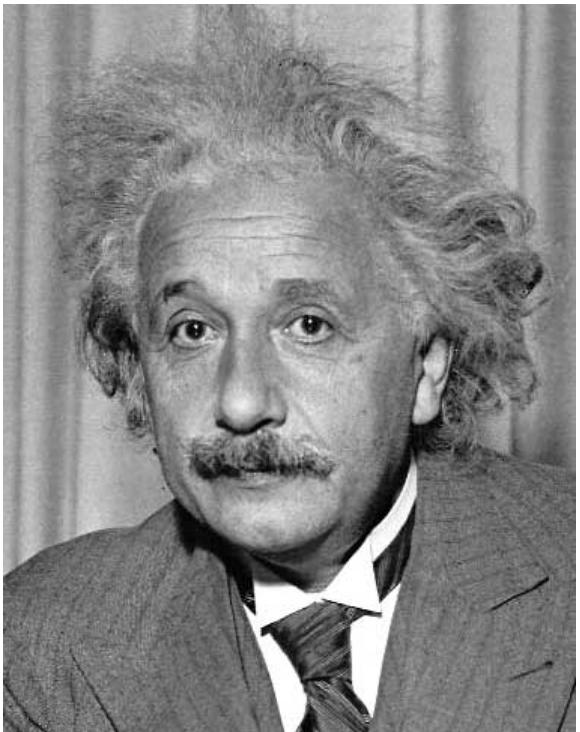


before



after

Other filters



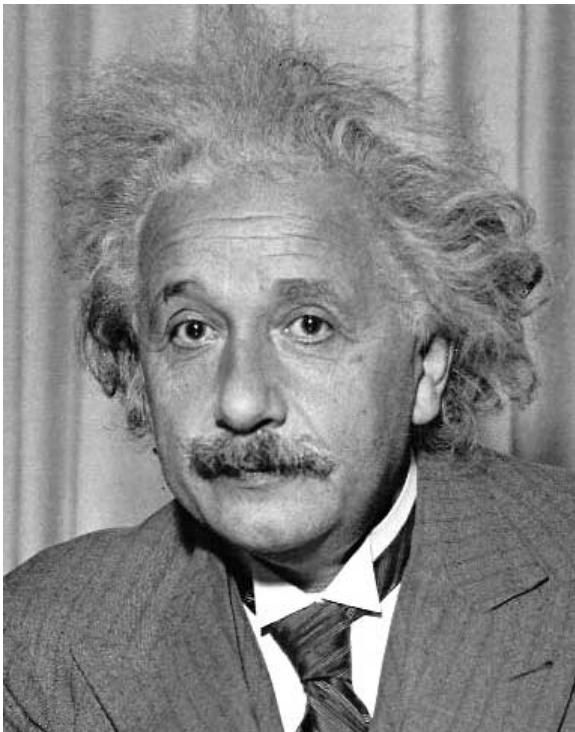
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Other filters



1	2	1
0	0	0
-1	2	-1

Sobel



Horizontal Edge
(absolute value)

Basic gradient filters

Horizontal Gradient

0	0	0
-1	0	1
0	0	0

or

-1	0	1
----	---	---

Vertical Gradient

0	-1	0
0	0	0
0	1	0

or

-1
0
1

Questions

Write as filtering operations, plus some pointwise operations: +, -, .*, >, etc.

1. Sum of four adjacent neighbors plus 1
2. Sum of squared values of 3x3 windows around each pixel:
3. Center pixel value is larger than the average of the pixel values to the left and right:

More properties

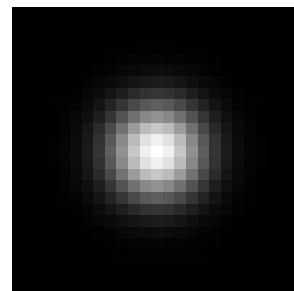
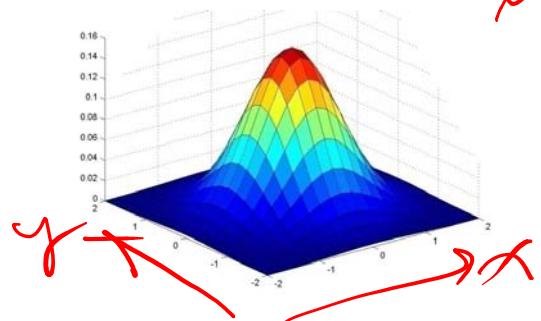
- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$
- Identity: unit impulse $e = [0, 0, 1, 0, 0]$, $a * e = a$

Gaussian Filter

$$\underline{x} \in \mathbb{R}^2$$

$$\exp\left(-\frac{\|x\|^2}{2\sigma^2}\right)$$

- Spatially-weighted averaging

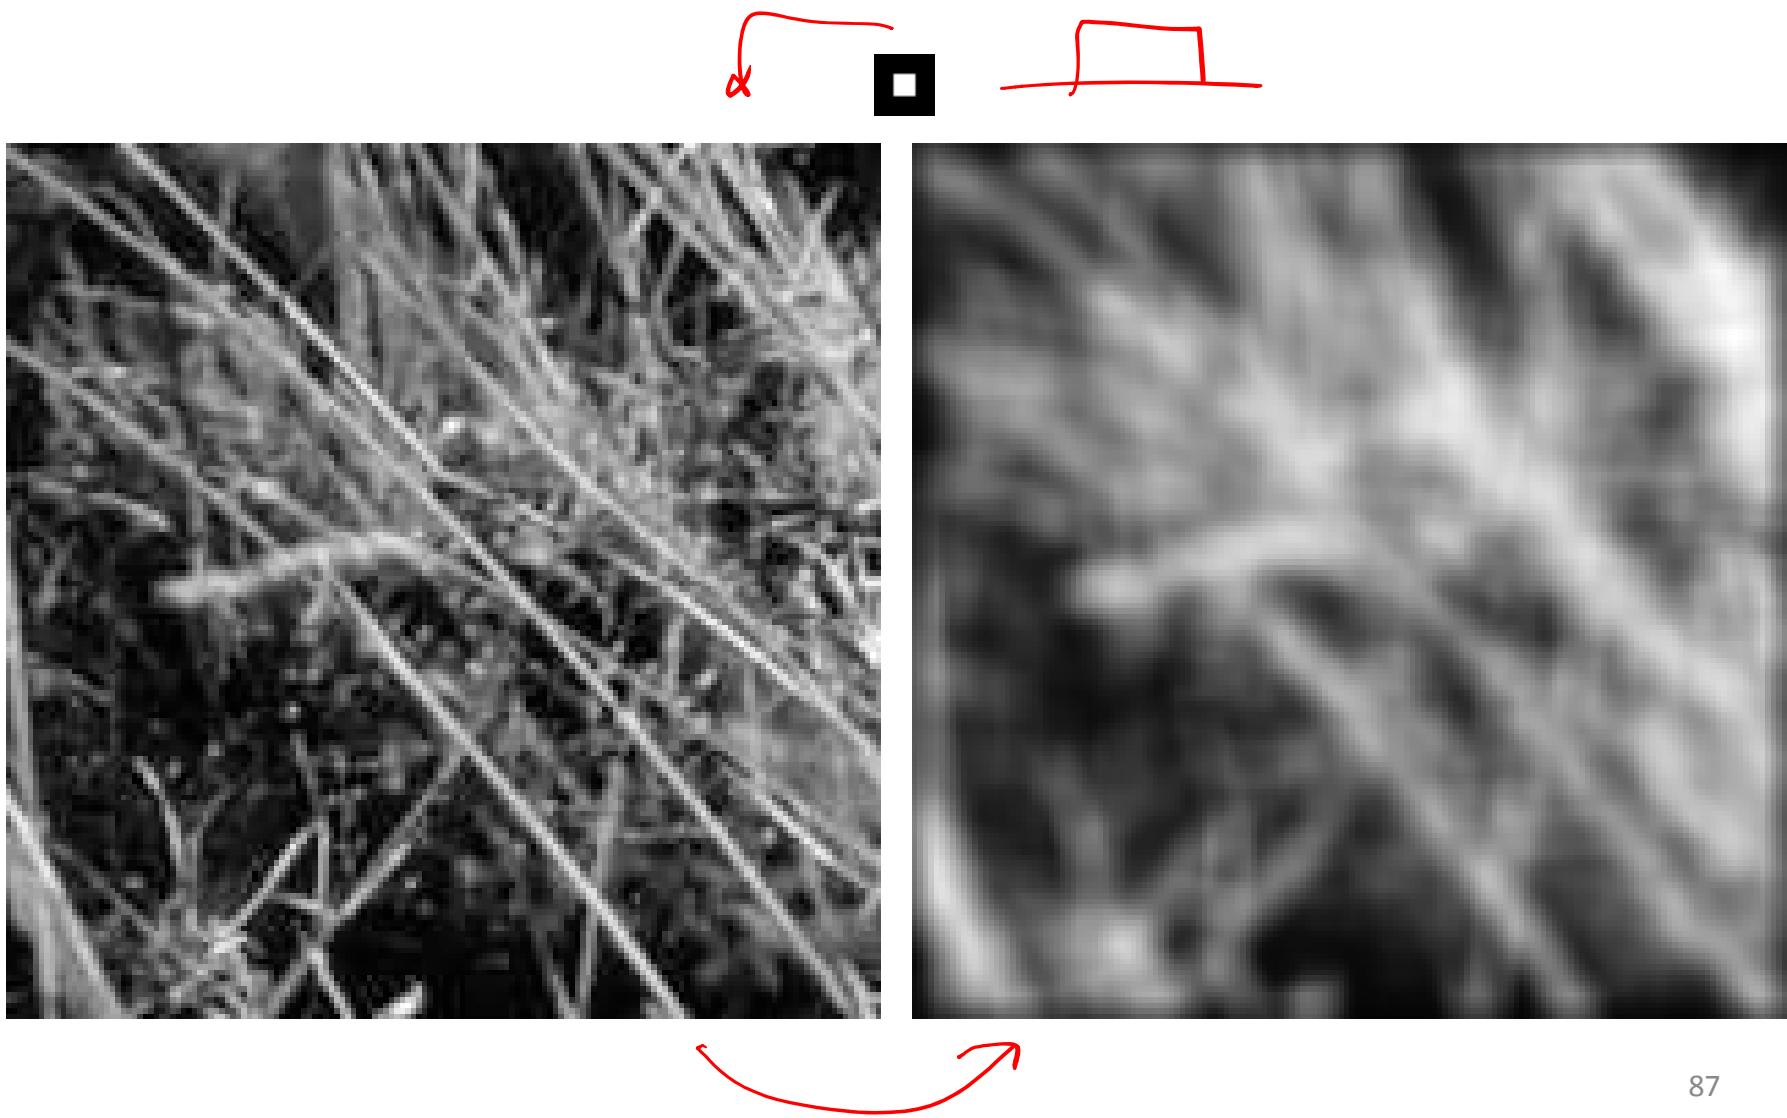


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

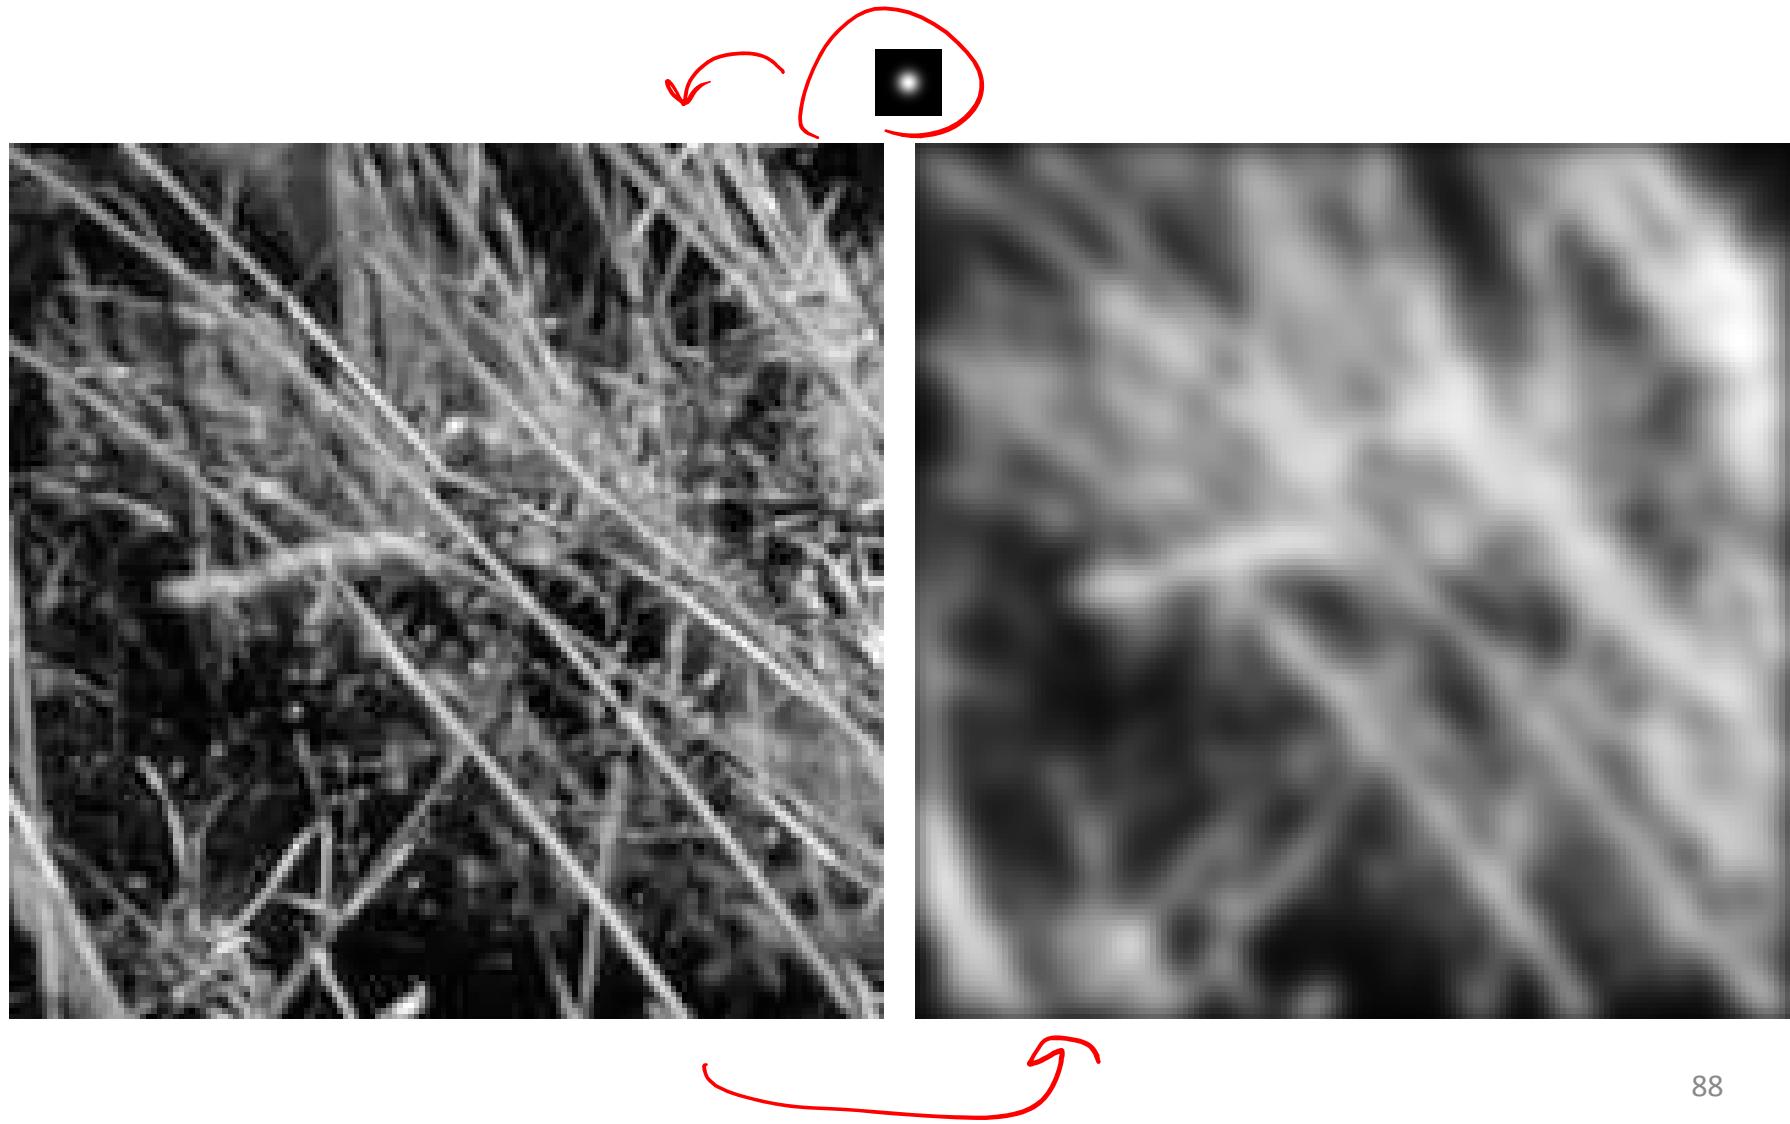
5 x 5, $\sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

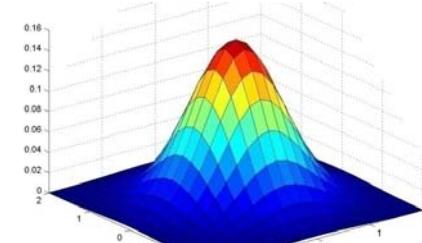
Smoothing with a Box Filter



Smoothing with a Gaussian Filter



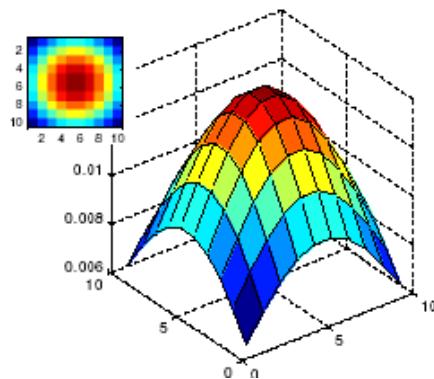
Gaussian Filter



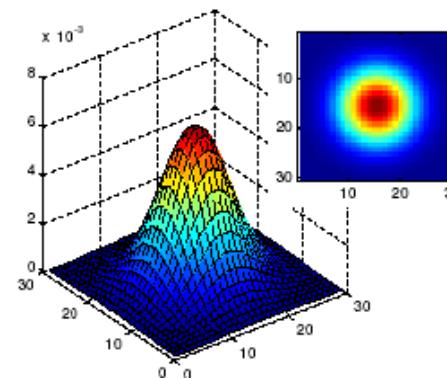
- Remove “high-frequency” components from the image (i.e., low-pass filter)
 - Images become more smooth
- Convolution with self is another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernels
 - Factors into product of two 1D Gaussians

Gaussian Filter

- What parameters matter here?
- **Size** of filter (or kernel/mask)



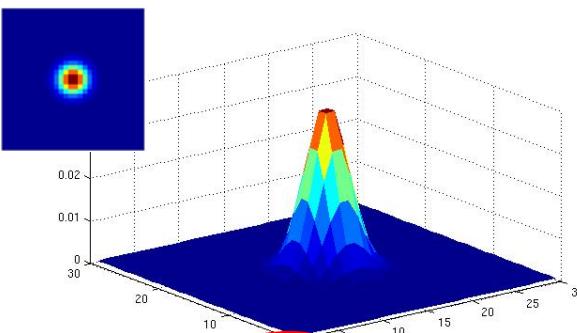
$\sigma = 5$ with 10×10 kernel



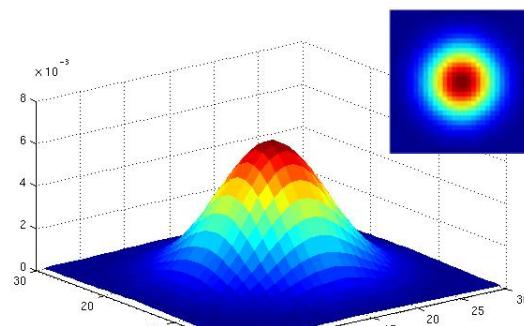
$\sigma = 5$ with 30×30 kernel

Gaussian Filter

- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing



$\sigma = 2$
with 30 x 30 kernel



$\sigma = 5$
with 30 x 30 kernel

Separability of the Gaussian Filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

$\underline{x} \in \mathbb{R}^d$

$\boxed{\frac{1}{(\underline{x}^\top \Sigma^{-1} \underline{x})^{\frac{d}{2}} \cdot \sigma^d}} \cdot \exp\left(-\frac{\|\underline{x}\|^2}{2\sigma^2}\right)$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

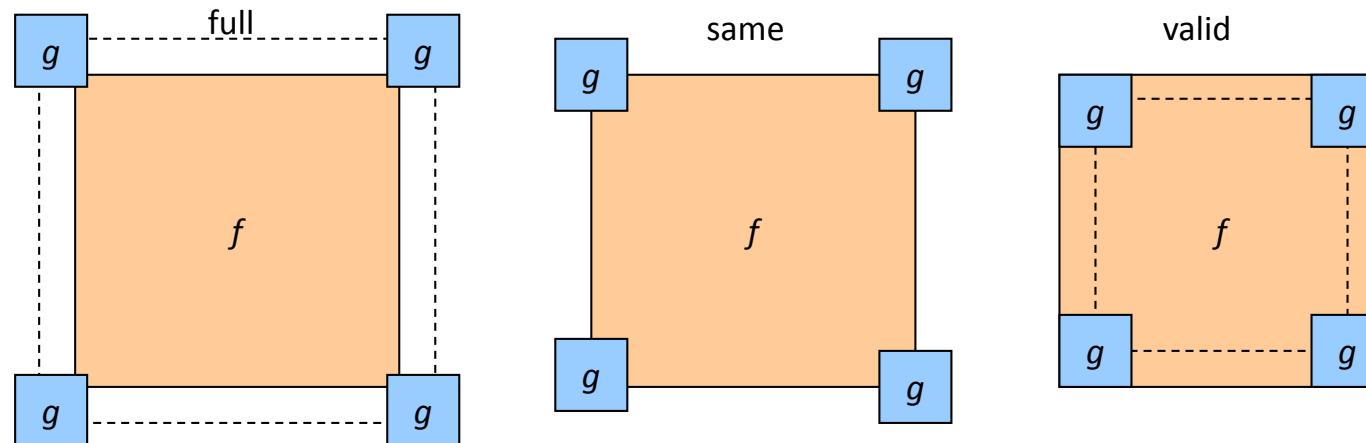
In this case, the two functions are the (identical) 1D Gaussian

Separability of the Gaussian Filter

- Why is separability useful in practice?
- Separability means that a 2D convolution can be reduced to two 1D convolutions (one among rows and one among columns)
- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$
- What if the kernel is separable?
 - $O(n^2 m)$

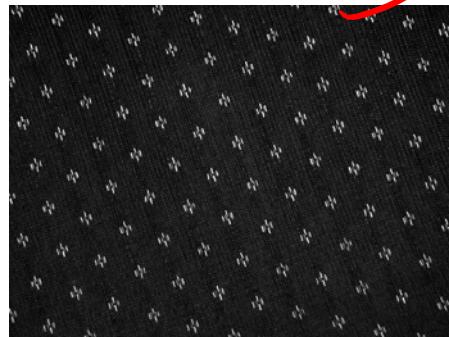
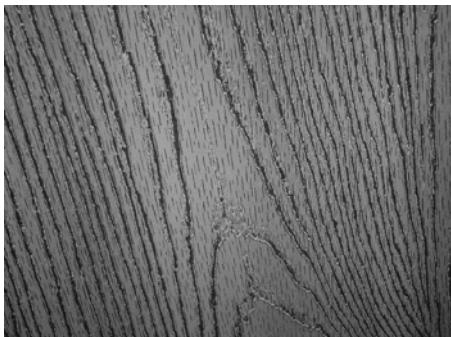
Practical matters

- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
 - *shape* = ‘full’: output size is sum of sizes of *f* and *g*
 - *shape* = ‘same’: output size is same as *f*
 - *shape* = ‘valid’: output size is difference of sizes of *f* and *g*



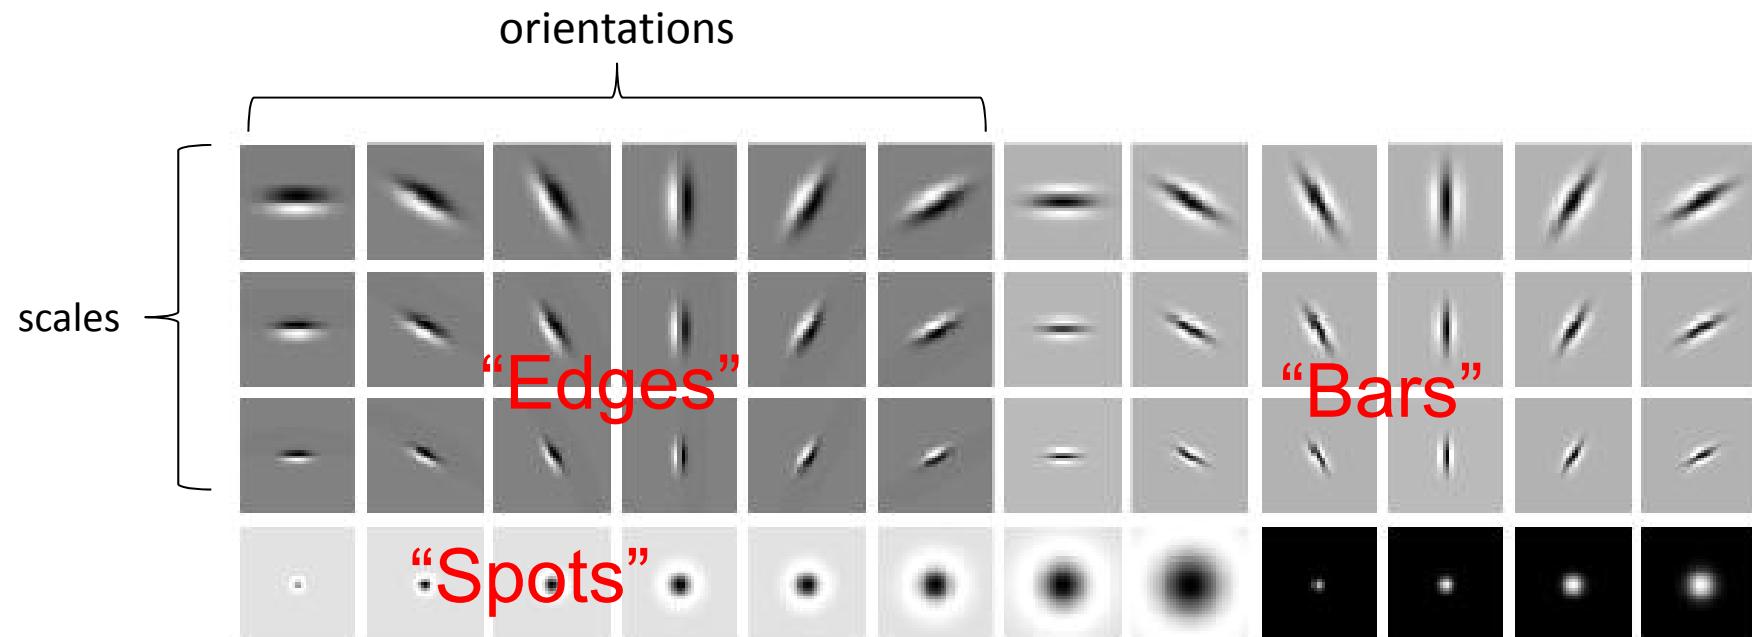
Applications: Representing Image Texture

- Texture: Orientation, Material, and Scale



How to Describe Textures?

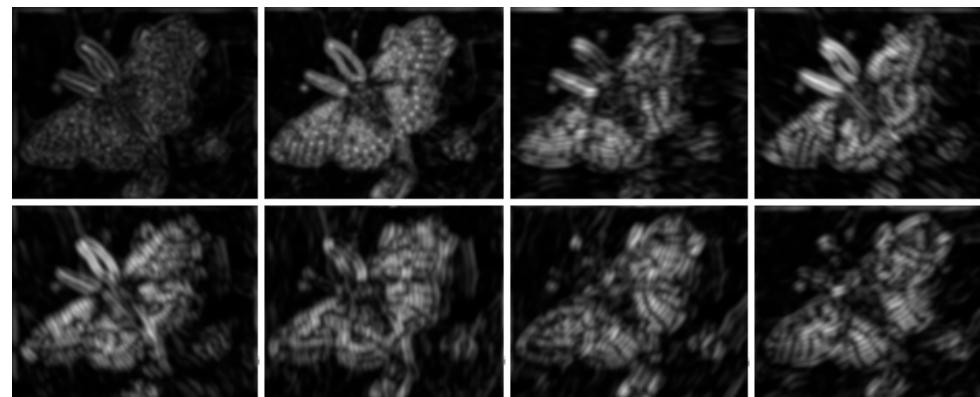
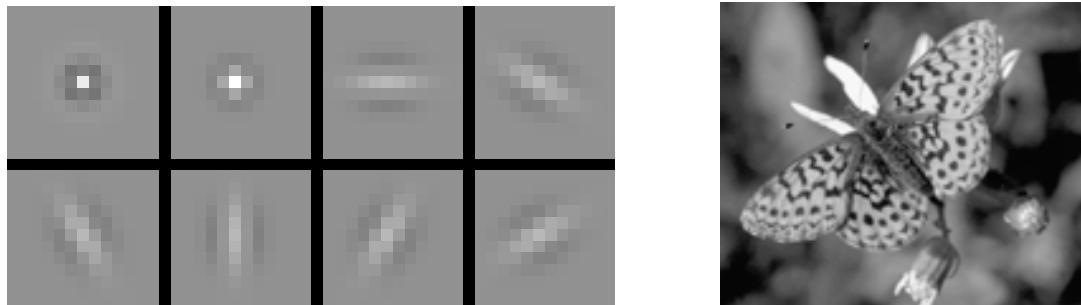
- Filter banks
 - Typically overcomplete representation
 - E.g., wavelet or more complex filters (basis functions)



Code for filter banks: www.robots.ox.ac.uk/~vqq/research/texclass/filters.html

How to Describe Textures?

- Filter banks
 - Typically overcomplete representation
 - E.g., wavelet or more complex filters (basis functions)
 - Example



Can you match the filtering response?

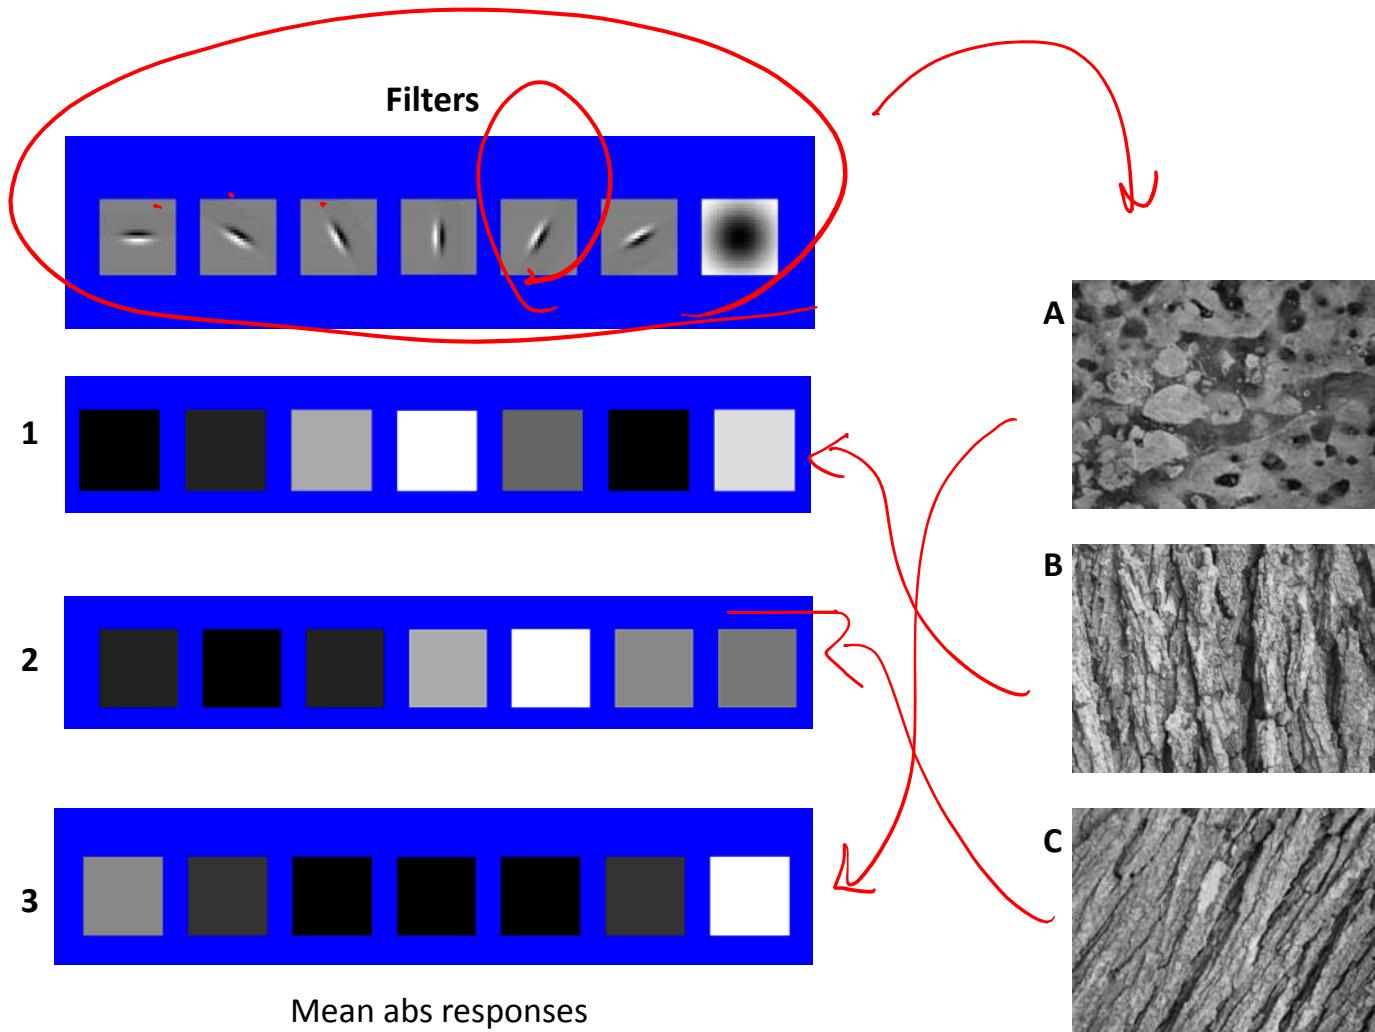


Image Denoising with Various Noise Types

- ✓
 - Salt and pepper
 - Contains random occurrences of black/white pixels
 - Impulse noise
 - Contains random occurrences of white pixels
 - Gaussian noise
 - Random pixels with variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise

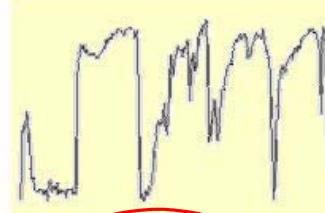
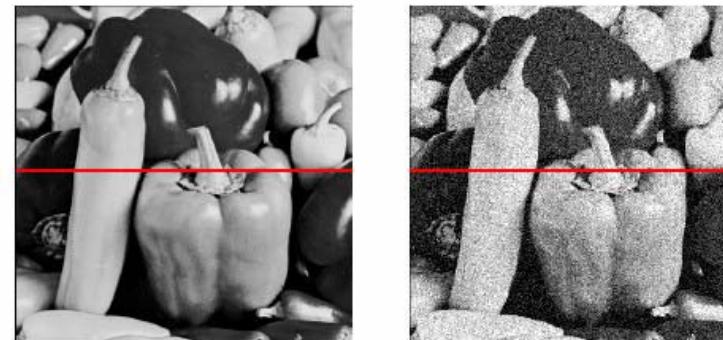


Gaussian noise

Gaussian Noise

- Assumption: independent, zero-mean noise
- Mathematical model: sum of many independent factors
- Good for small standard deviations

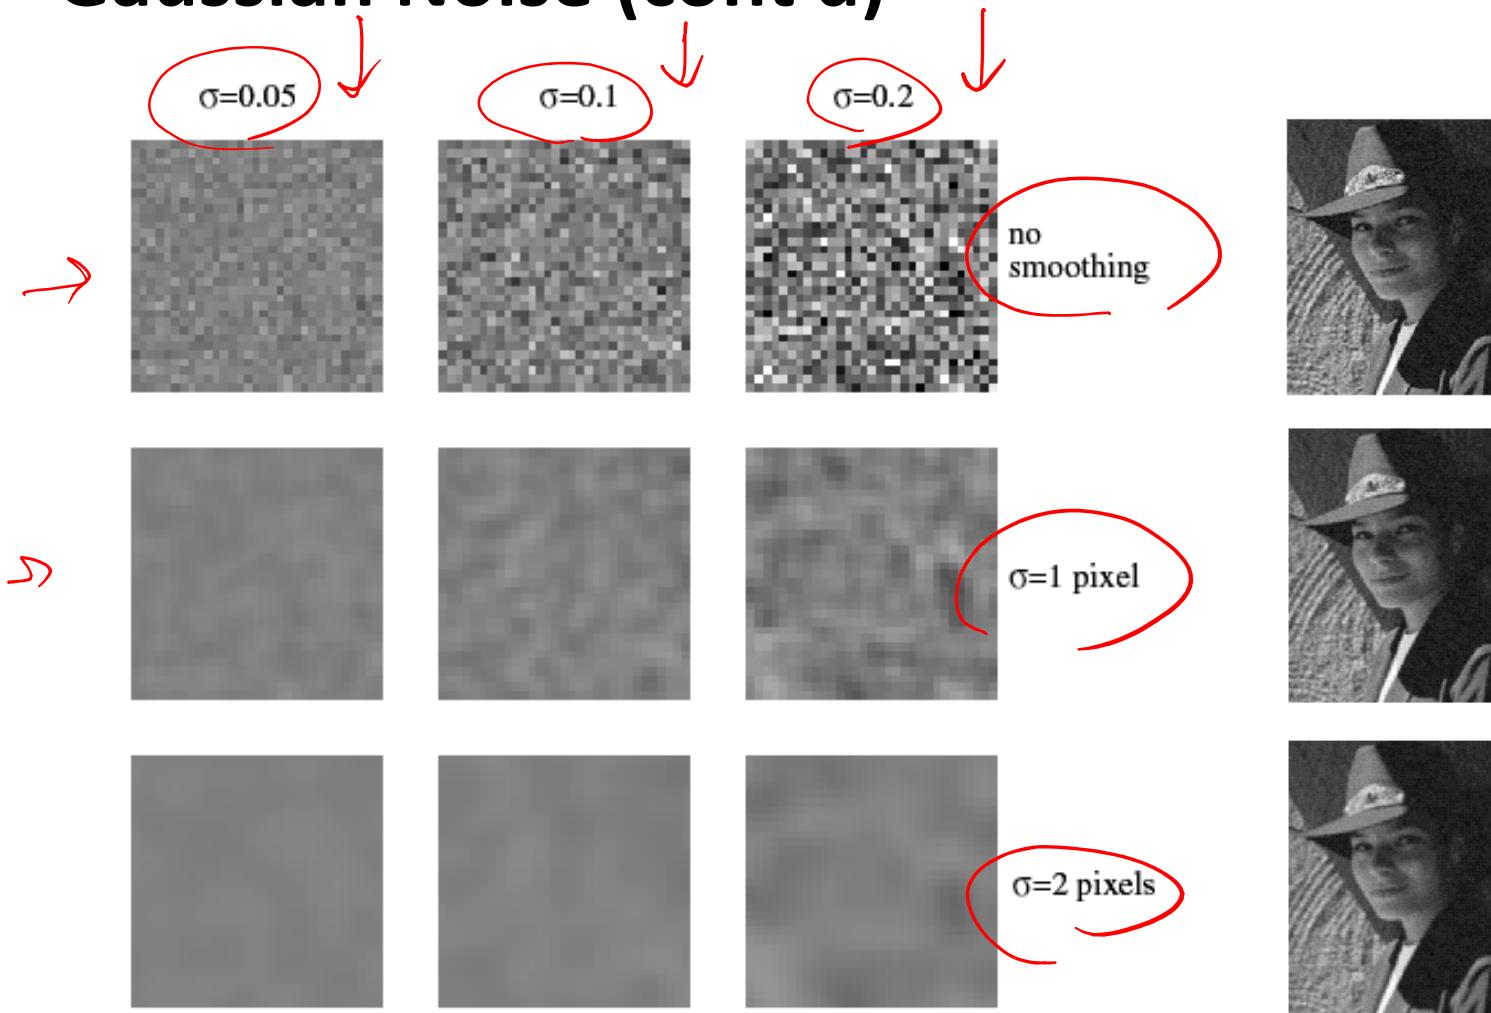
Image
Noise



$$f(x, y) = \widehat{f}(x, y) + \eta(x, y)$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

Gaussian Noise (cont'd)



- Smoothing with larger σ suppresses noise but also results in blurry images.

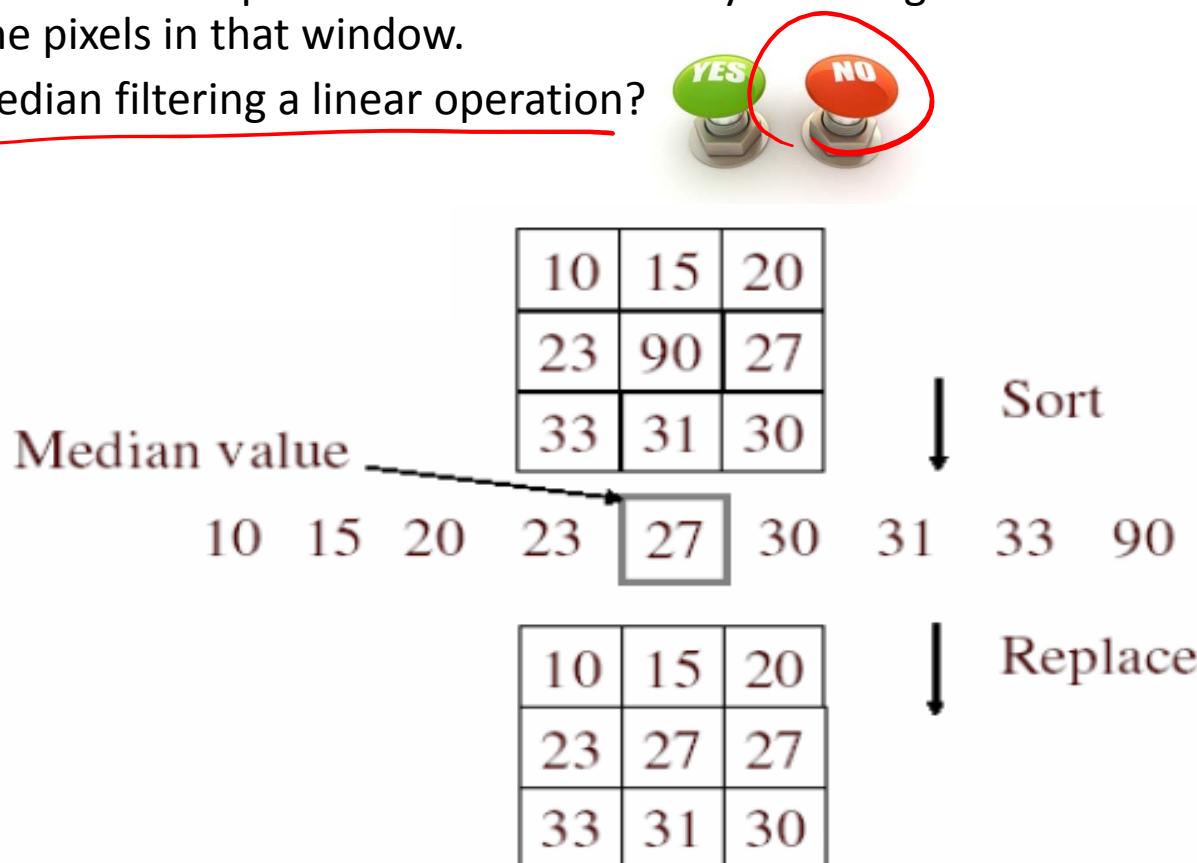
Salt-and-Pepper Noise

- What's wrong with the results?



Alternative Idea...

- Median Filtering
 - A median filter operates over a window by selecting the median intensity of the pixels in that window.
 - Is median filtering a linear operation?

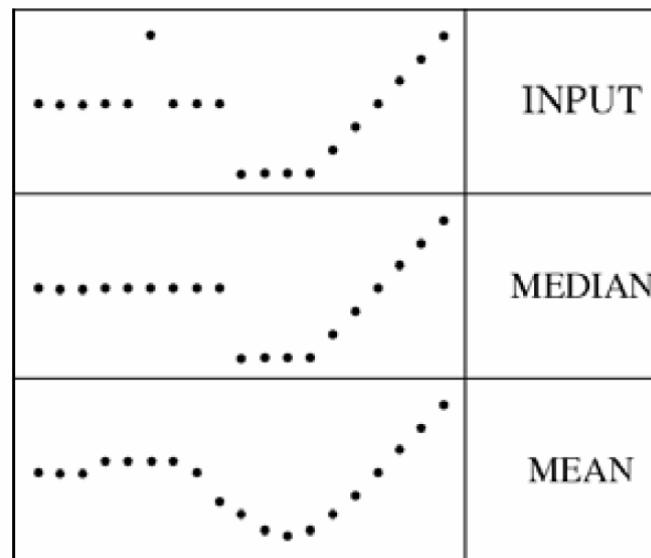


Alternative Idea...

- Median Filtering
 - A median filter operates over a window by selecting the median intensity of the pixels in that window.
 - Is median filtering a linear operation?
 - What advantage does median filter have over the Gaussian one?
 - Robustness to **outliers**.

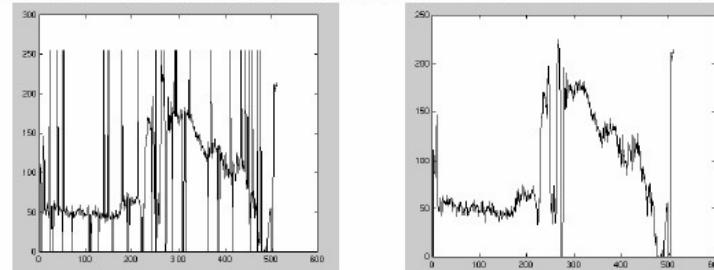
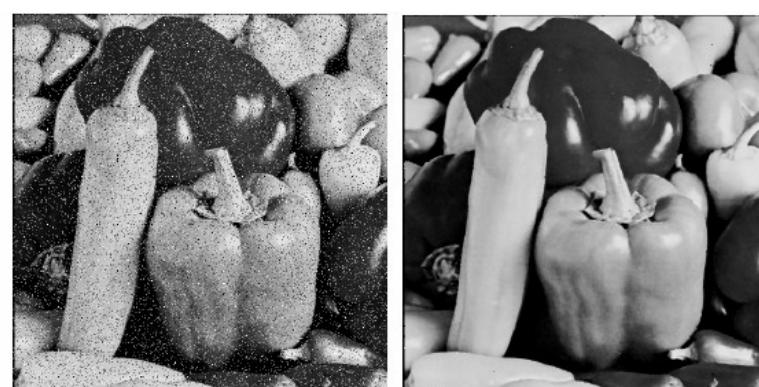


filters have width 5 :



Alternative Idea...

- Median Filtering
 - A median filter operates over a window by selecting the median intensity of the pixels in that window.
 - What advantage does median filter have over the Gaussian one?
 - Robustness to outliers.
 - Example [Salt-and-pepper noise](#) [Median filtered](#)



Other Non-Linear Filters

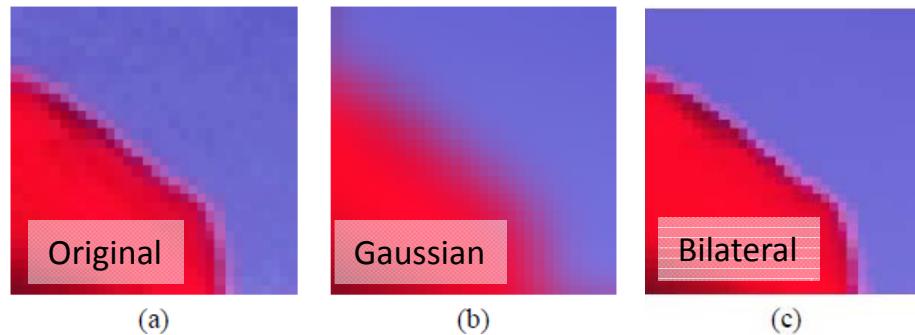
- Weighted median (pixels further from center count less)
- Clipped mean (average, ignoring few brightest and darkest pixels)
- Bilateral filtering (weight by spatial distance *and* intensity difference)



Bilateral filtering

Bilateral Filters

- Edge preserving: weights similar pixels more



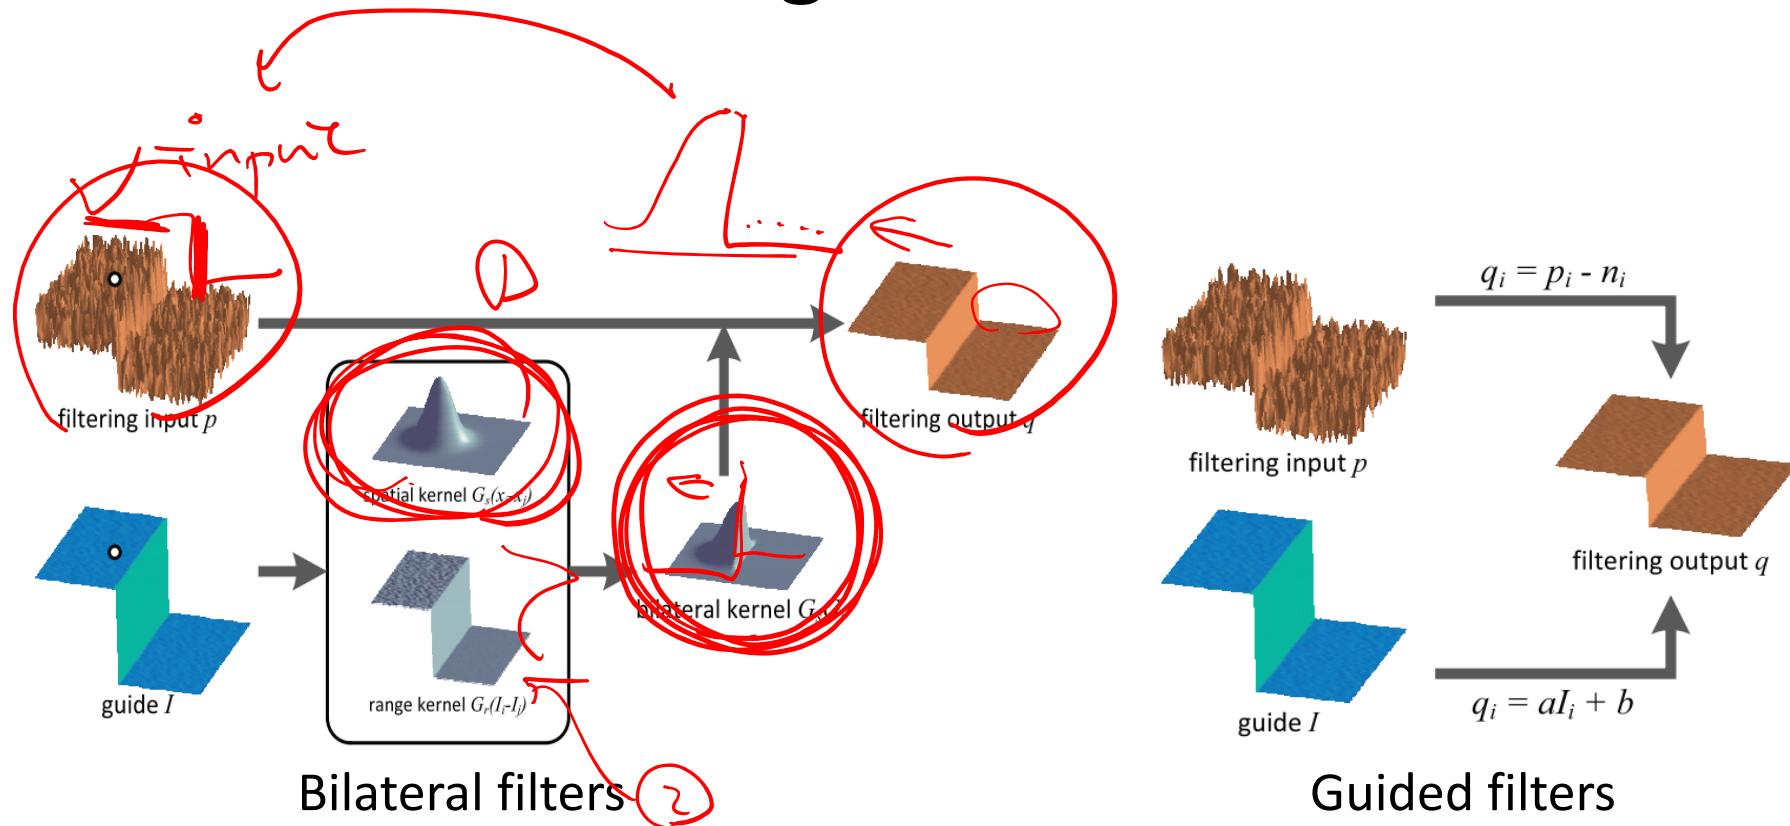
$$I_p^b = \frac{1}{W_p^b} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

with $W_p^b = \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|)$

spatial similarity (e.g., intensity)

(Handwritten annotations: 'spatial' and 'similarity (e.g., intensity)' are circled in red, with arrows pointing to their respective terms in the equation. A small red bell-shaped curve is drawn below the first summation term.)

BF as Guided Image Filters



`B = imguidedfilter(A, G);`